

Assembly Programming: Data

Today

■ Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

■ Structures

- Allocation
- Access
- Alignment


Array Allocation

■ Basic Principle

$T \ A[L];$

- Array of data type T and length L
- Contiguously allocated region of $L * \text{sizeof}(T)$ bytes in memory

12 bytes all adjacent




Array Access

Basic Principle

T $A[L];$

- Array of data type T and length L
- Identifier A can be used as a pointer to array element 0: Type T^*



Reference

Type

Value

<code>val[4]</code>	<u>int</u>	3
<code>val</code>	int *	x
<code>val+1</code>	int *	$x+4$
<code>&val[2]</code>	int *	$x+8$
<code>val[5]</code> ← crash or garbage data	int	??
<code>*(val+1)</code> ← pointer arithmetic dereference	int	5
<code>val + i</code>	int *	$x+4i$

$x+4i$

Array Example


```
#define ZLEN 5
typedef int zip_dig[ZLEN];

zip_dig cmu = { 1, 5, 2, 1, 3 };
zip_dig mit = { 0, 2, 1, 3, 9 };
zip_dig ucb = { 9, 4, 7, 2, 0 };
```


20 byte blocks

20 byte blocks


zip_dig cmu;



zip_dig mit;



zip_dig ucb;




Cmu may/may not end and forth with 2 m

- Declaration "zip_dig cmu" equivalent to "int cmu[5]"
- Example arrays were allocated in successive 20 byte blocks
- Not guaranteed to happen in general

Two Sigma

Array Accessing Example

zip_dig cmu;



```
int get_digit
(zip_dig z, int digit)
{
    return z[digit];
}
```

Assembly

```
# %rdi = z
# %rsi = digit
movl (%rdi,%rsi,4), %eax # z[digit]
```

Handwritten notes:
* $4(rsi) + rdi$
put into rax
return value
desired

- Register `%rdi` contains starting address of array
- Register `%rsi` contains array index
- Desired digit at `%rdi + 4*%rsi`
- Use memory reference `(%rdi,%rsi,4)`

Array Loop Example

```
void zincr(zip_dig z) {  
    size_t i;  
    for (i = 0; i < ZLEN; i++)  
        z[i]++;  
}
```

cf xrax + rdi

```
# %rdi = z  
movl    $0, %eax           # i = 0  
jmp     .L3                # goto middle  
.L4:                        # loop:  
addl    $1, (%rdi,%rax,4)  # z[i]++  
addq    $1, %rax           # i++  
.L3:                        # middle  
cmpq    $4, %rax           # i:4  
jbe     .L4                # if <=, goto loop  
ret
```

we have to increment it

unsigned equal

Multidimensional (Nested) Arrays

Declaration

$T \ A[R][C];$


- 2D array of data type T
- R rows, C columns
- Type T element requires K bytes

Array Size

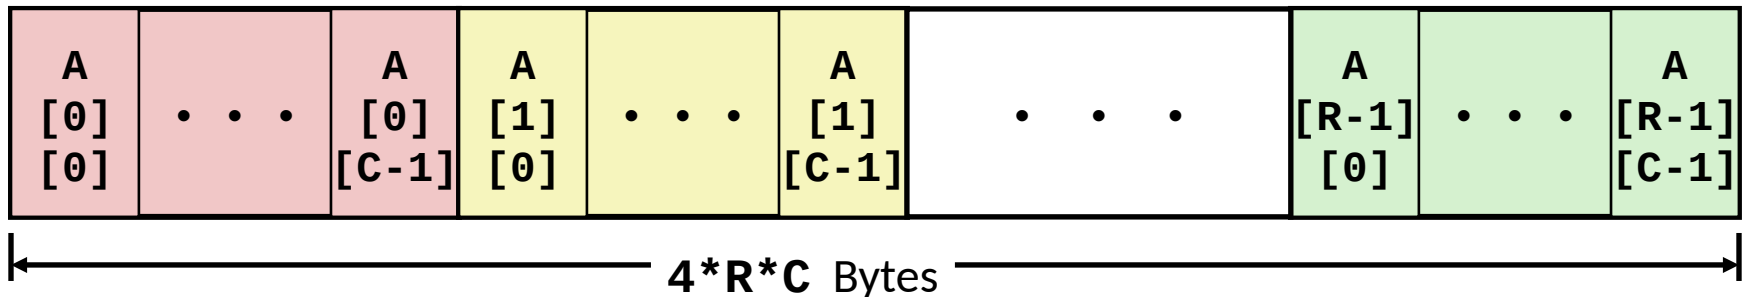
- $R * C * K$ bytes

Arrangement

- Row-Major Ordering




`int A[R][C];`



Type like int so integers L/R C

Nested Array Example

```
#define PCOUNT 4
zip_dig pgh[PCOUNT] =
    {{1, 5, 2, 0, 6},
     {1, 5, 2, 1, 3 },
     {1, 5, 2, 1, 7 },
     {1, 5, 2, 2, 1 }};
```




- “**zip_dig pgh[4]**” equivalent to “**int pgh[4][5]**”
 - Variable **pgh**: array of 4 elements, allocated contiguously
 - Each element is an array of 5 **int**’s, allocated contiguously
- “**Row-Major**” ordering of all elements in memory

Nested Array Row Access


■ Row Vectors

- $A[i]$ is array of C elements
- Each element of type T requires K bytes
- Starting address $A + i * (C * K)$

`int A[R][C];`



Nested Array Row Access Code



$pgh + i \cdot C$

```
int *get_pgh_zip(int index)
{
    return pgh[index];
}
```

return a pointer to the index

```
# %rdi = index
leaq (%rdi,%rdi,4),%rax # 5 * index
leaq pgh(,%rax,4),%rax # pgh + (20 * index)
```

$4(rdi) + rdi$

■ Row Vector

- `pgh[index]` is array of 5 `int`'s
- Starting address `pgh+20*index`

■ Machine Code


- Computes and returns address
- Compute as `pgh + 4*(index+4*index)`

Nested Array Element Access


■ Array Elements

- $A[i][j]$ is element of type T , which requires K bytes
- Address $A + i * (C * K) + j * K = A + (i * C + j) * K$ ← formula

`int A[R][C];`



Nested Array Element Access Code



```
int get_pgh_digit
(int index, int dig)
{
    return pgh[index][dig];
}
```

$$A + (ic + j) \cdot K$$

Handwritten note: $pgh + (\text{index} \cdot 5 + \text{dig}) \cdot 4$. The '5' is circled and labeled '5x5 in each row'. The '4' is labeled '4x4 in each column'.

```
leaq    (%rdi,%rdi,4), %rax    # 5*index
addl    %rax, %rsi             # 5*index+dig
movl    pgh(,%rsi,4), %eax     # M[pgh + 4*(5*index+dig)]
```

■ Array Elements


- `pgh[index][dig]` is `int`
- Address: `pgh + 20*index + 4*dig`
 - `= pgh + 4*(5*index + dig)`

Multi-Level Array Example

```
zip_dig cmu = { 1, 5, 2, 1, 3 };  
zip_dig mit = { 0, 2, 1, 3, 9 };  
zip_dig ucb = { 9, 4, 7, 2, 0 };
```


```
#define UCOUNT 3  
int *univ[UCOUNT] = {mit, cmu, ucb};
```

- Variable `univ` denotes array of 3 elements
- Each element is a pointer
 - 8 bytes
- Each pointer points to array of `int`'s



Element Access in Multi-Level Array

```
int get_univ_digit
(size_t index, size_t digit)
{
    return univ[index][digit];
}
```



```
salq    $2, %rsi          # 4*digit
addq    univ(,%rdi,8), %rsi # p = univ[index] + 4*digit
movl    (%rsi), %eax       # return *p
ret
```

■ Computation

- Element access **Mem[Mem[univ+8*index]+4*digit]**
- Must do two memory reads
 - First get pointer to row array
 - Then access element within array

Handwritten notes:
 univ + (8 * index) → pointer to row array
 + ((univ + 8 * index) + digit * 4)


Array Element Accesses

Nested array


```
int get_pgh_digit
(size_t index, size_t digit)
{
    return pgh[index][digit];
}
```

Multi-level array

```
int get_univ_digit
(size_t index, size_t digit)
{
    return univ[index][digit];
}
```



more compact in memory



Accesses looks similar in C, but address computations very different:

$\text{Mem}[\text{pgh} + 20 * \text{index} + 4 * \text{digit}]$

$\text{Mem}[\text{Mem}[\text{univ} + 8 * \text{index}] + 4 * \text{digit}]$

Today

■ Arrays


- One-dimensional
- Multi-dimensional (nested)
- Multi-level

■ Structures

- Allocation
- Access
- Alignment

Structure Representation


```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



- **Structure represented as block of memory**
 - Big enough to hold all of the fields
- **Fields ordered according to declaration**
 - Even if another ordering could yield a more compact representation
- **Compiler determines overall size + positions of fields**
 - Machine-level program has no understanding of the structures in the source code

Generating Pointer to Structure Member

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



■ Generating Pointer to Array Element

- Offset of each structure member determined at compile time
- Compute as $r + 4 \cdot idx$

```
int *get_ap  
(struct rec *r, size_t idx)  
{  
    return &r->a[idx];  
}
```


```
# r in %rdi, idx in %rsi  
leaq (%rdi,%rsi,4), %rax  
ret
```

Following Linked List

■ C Code

```
void set_val
(struct rec *r, int val)
{
    while (r) {
        int i = r->i;
        r->a[i] = val;
        r = r->next;
    }
}
```

```
struct rec {
    int a[4];
    int i;
    struct rec *next;
};
```




Element i

Register	Value
%rdi	r
%rsi	val

```
.L11:                                # loop:
    movslq    16(%rdi), %rax          # i = M[r+16]
    movl      %esi, (%rdi,%rax,4)     # M[r+4*i] = val
    movq      24(%rdi), %rdi         # r = M[r+24]
    testq     %rdi, %rdi             # Test r
    jne       .L11                  # if !=0 goto loop
```

Structures & Alignment


■ Unaligned Data



```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

■ Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K



Specific Cases of Alignment (x86-64)

- 1 byte: **char**, ...

- no restrictions on address

- 2 bytes: **short**, ...

- lowest 1 bit of address must be 0_2

- 4 bytes: **int**, **float**, ...


- lowest 2 bits of address must be 00_2

- 8 bytes: **double**, **long**, **char ***, ...

- lowest 3 bits of address must be 000_2

- 16 bytes: **long double** (GCC on Linux)

- lowest 4 bits of address must be 0000_2



Satisfying Alignment with Structures

■ Within structure:

- Must satisfy each element's alignment requirement


■ Overall structure placement

- Each structure has alignment requirement K
 - K = Largest alignment of any element
- Initial address & structure length must be multiples of K

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

■ Example:


- K = 8, due to double element



Meeting Overall Alignment Requirement

- For largest alignment requirement K
- Overall structure must be multiple of K

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```



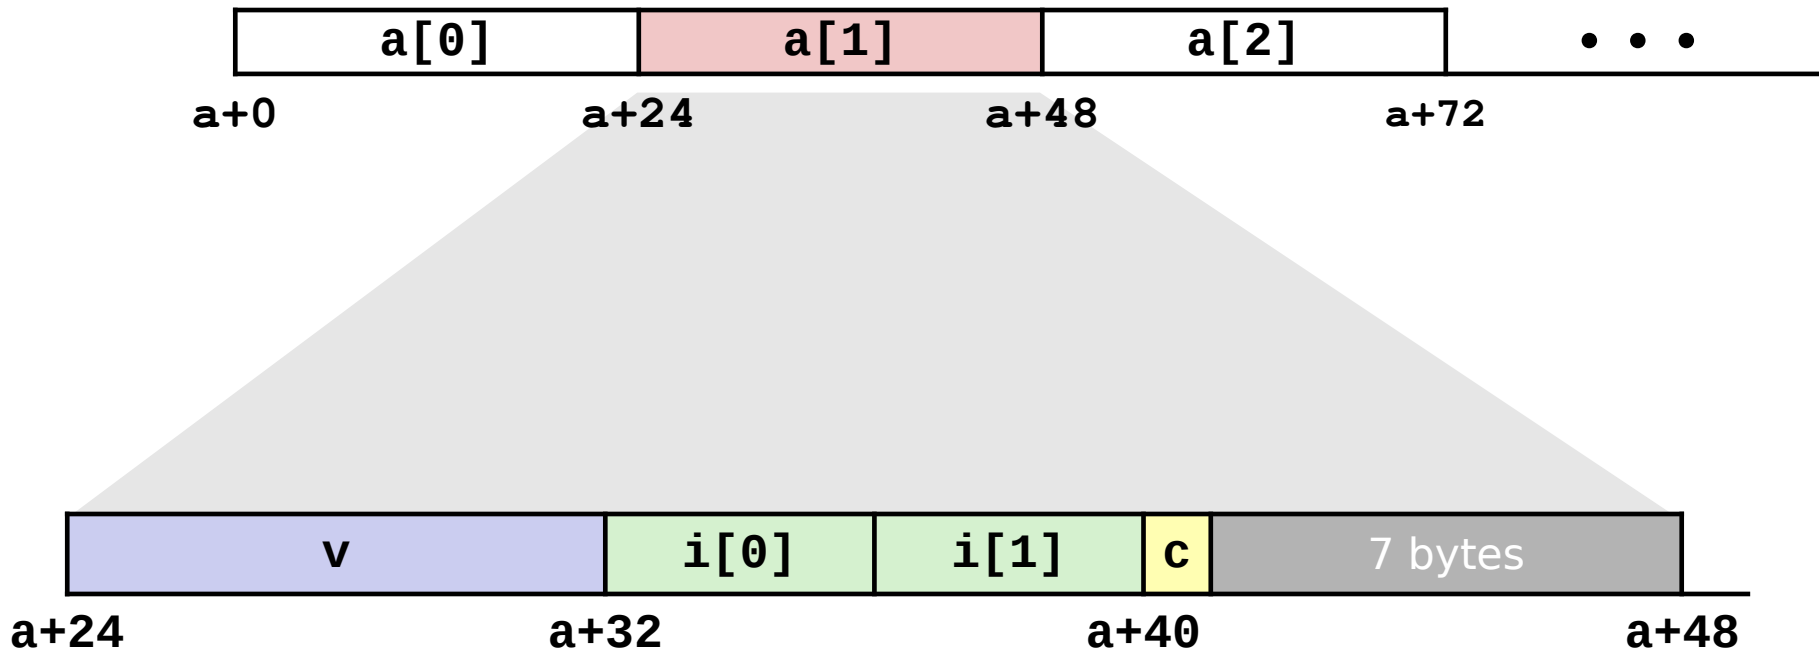
Multiple of $K=8$

An arrow points from the text 'Multiple of K=8' to the p+24 address mark on the diagram.

Arrays of Structures

- Overall structure length multiple of K
- Satisfy alignment requirement for every element


```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



Saving Space

- Put large data types first

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```

- Effect (K=4)

