

The Memory Hierarchy

Random-Access Memory (RAM)

■ Key features

- RAM is traditionally packaged as a chip.
- Basic storage unit is normally a cell (one bit per cell).
- Multiple RAM chips form a memory.

■ RAM comes in two varieties:

- SRAM (Static RAM)
- DRAM (Dynamic RAM)

SRAM vs DRAM Summary

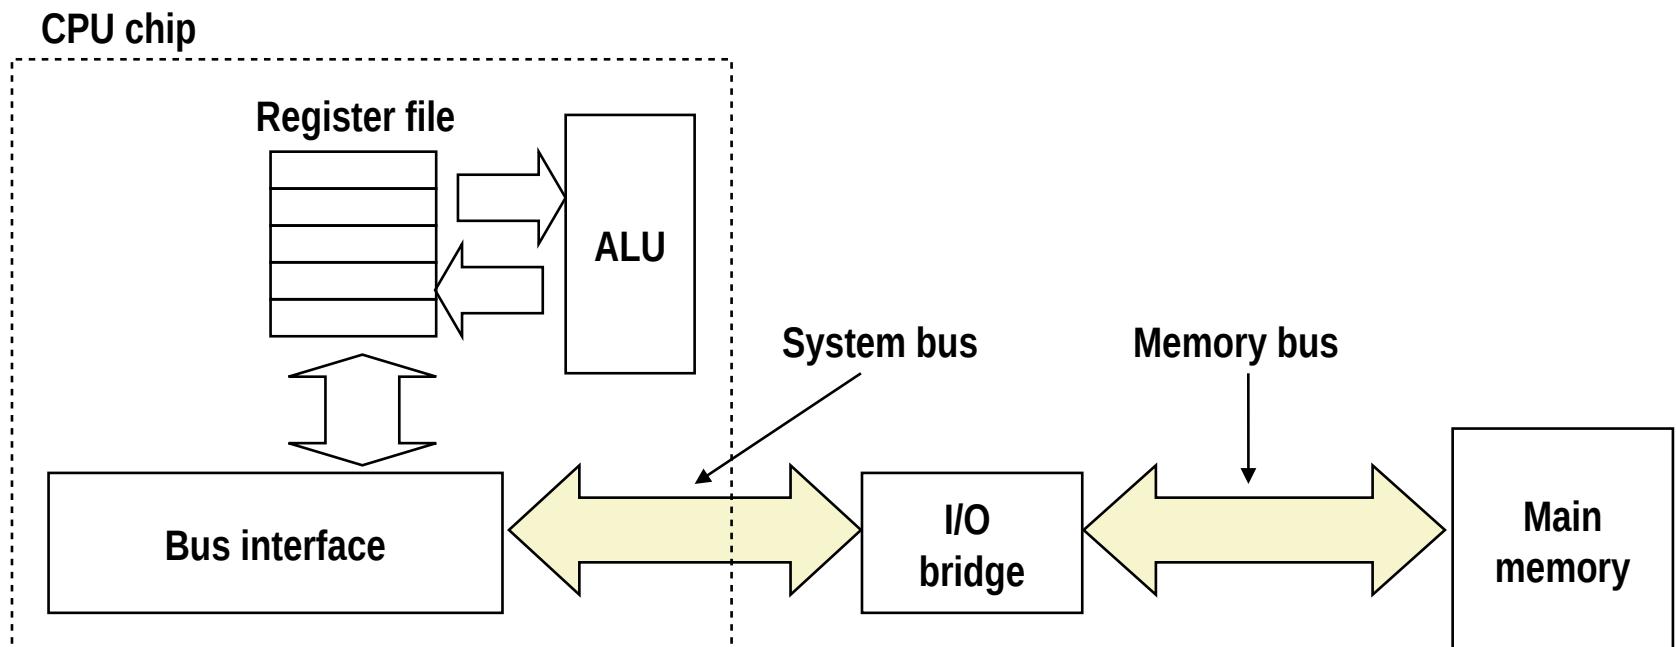
	Trans. per bit	Access time	Needs refresh?	Needs EDC?	Cost	Applications
SRAM	4 or 6	1X	No	Maybe	100x	Cache memories
DRAM	1	10X	Yes	Yes	1X	Main memories, frame buffers

Nonvolatile Memories

- DRAM and SRAM are volatile memories
 - Lose information if powered off.
 - Nonvolatile memories retain value even if powered off
 - Read-only memory (**ROM**): programmed during production
 - Programmable ROM (**PROM**): can be programmed once
 - Eraseable PROM (**EPROM**): can be bulk erased (UV, X-Ray)
 - Electrically eraseable PROM (**EEPROM**): electronic erase capability
 - Flash memory: EEPROMs. with partial (block-level) erase capability
 - Wears out after about 100,000 erasings
 - Uses for Nonvolatile Memories
 - Firmware programs stored in a ROM (BIOS, controllers for disks, network cards, graphics accelerators, security subsystems,...)
 - Solid state disks (replace rotating disks in thumb drives, smart phones, mp3 players, tablets, laptops,...)
 - Disk caches
- Stay
F firms*

Traditional Bus Structure Connecting CPU and Memory

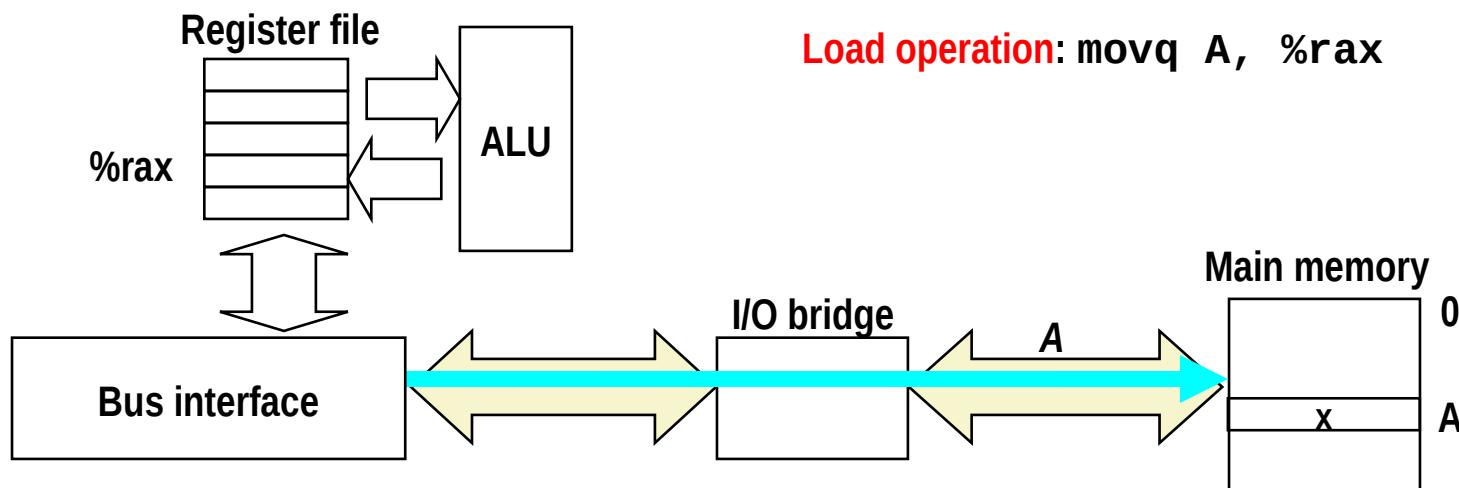
- A **bus** is a collection of parallel wires that carry address, data, and control signals.
- Buses are typically shared by multiple devices.



Place add reg A

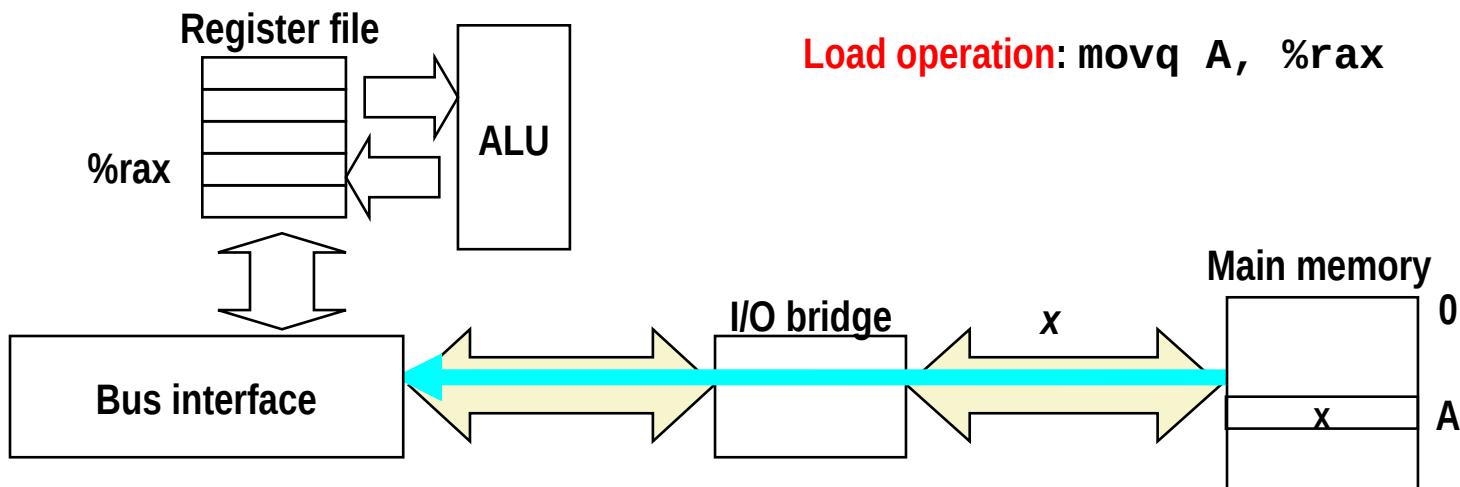
Memory Read Transaction (1)

- CPU places address A on the memory bus.



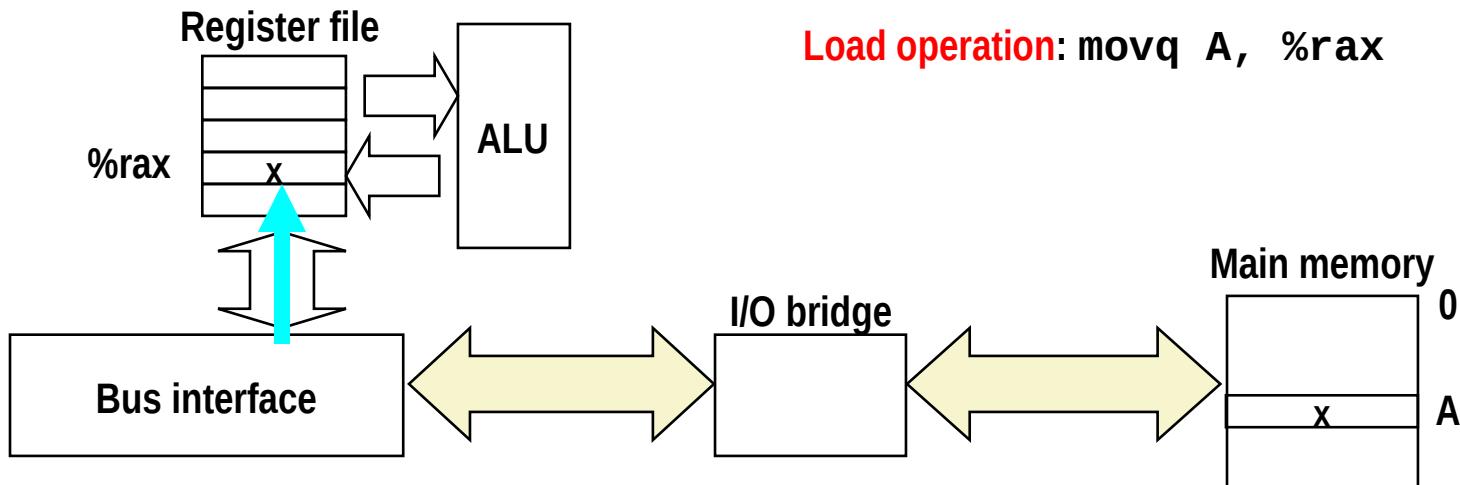
Memory Read Transaction (2)

- Main memory reads A from the memory bus, retrieves word x, and places it on the bus.



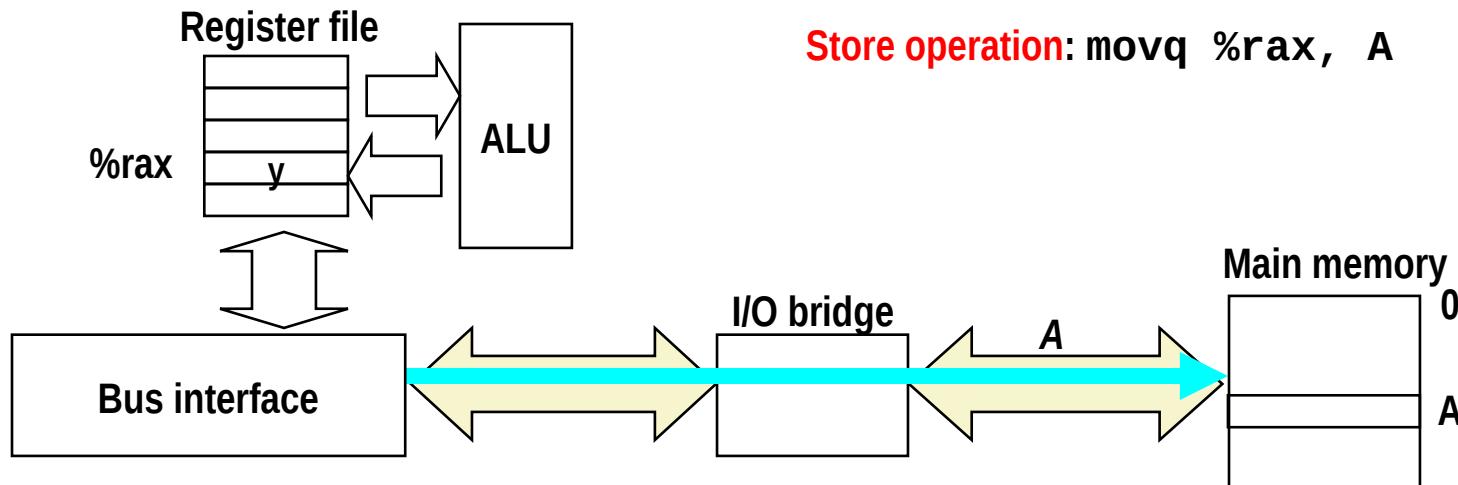
Memory Read Transaction (3)

- CPU read word x from the bus and copies it into register $\%rax$.



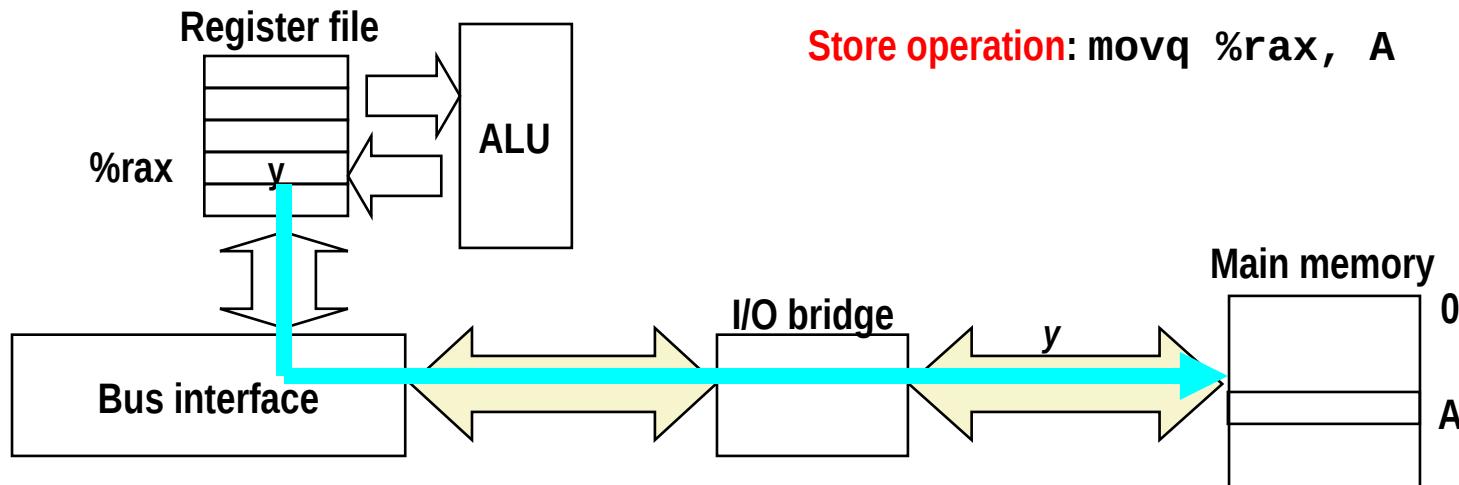
Memory Write Transaction (1)

- CPU places address A on bus. Main memory reads it and waits for the corresponding data word to arrive.



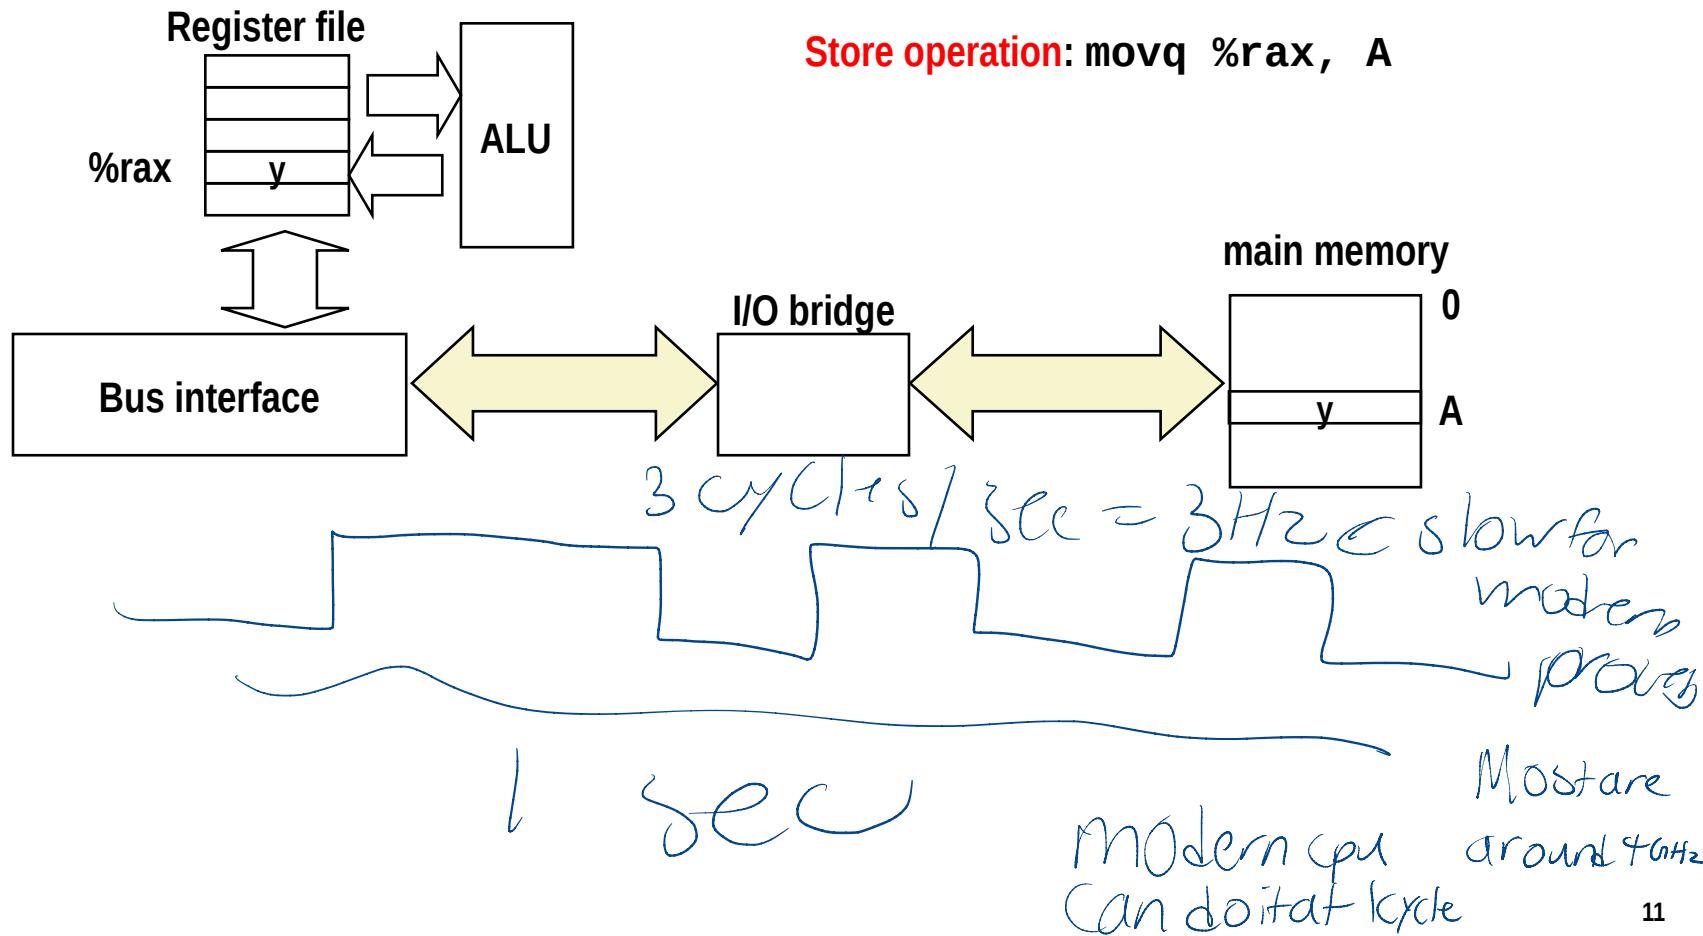
Memory Write Transaction (2)

- CPU places data word y on the bus.



Memory Write Transaction (3)

- Main memory reads data word y from the bus and stores it at address A.



Locality

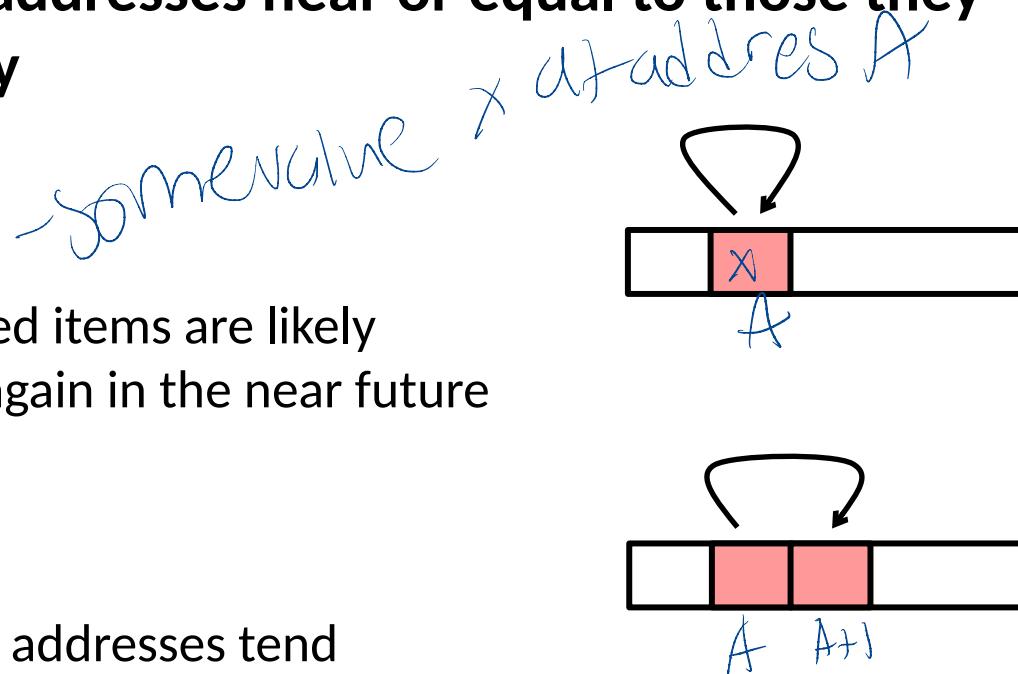
- **Principle of Locality:** Programs tend to use data and instructions with addresses near or equal to those they have used recently

- **Temporal locality:**

- Recently referenced items are likely to be referenced again in the near future

- **Spatial locality:**

- Items with nearby addresses tend to be referenced close together in time



Locality Example

load into CPU and execute

```
sum = 0;
for (i = 0; i < n; i++)
    sum += a[i];
return sum;
```

a[0]
a[1]
a[2]

■ Data references

- Reference array elements in succession (stride-1 reference pattern).
- Reference variable sum each iteration.

Spatial locality

Temporal locality

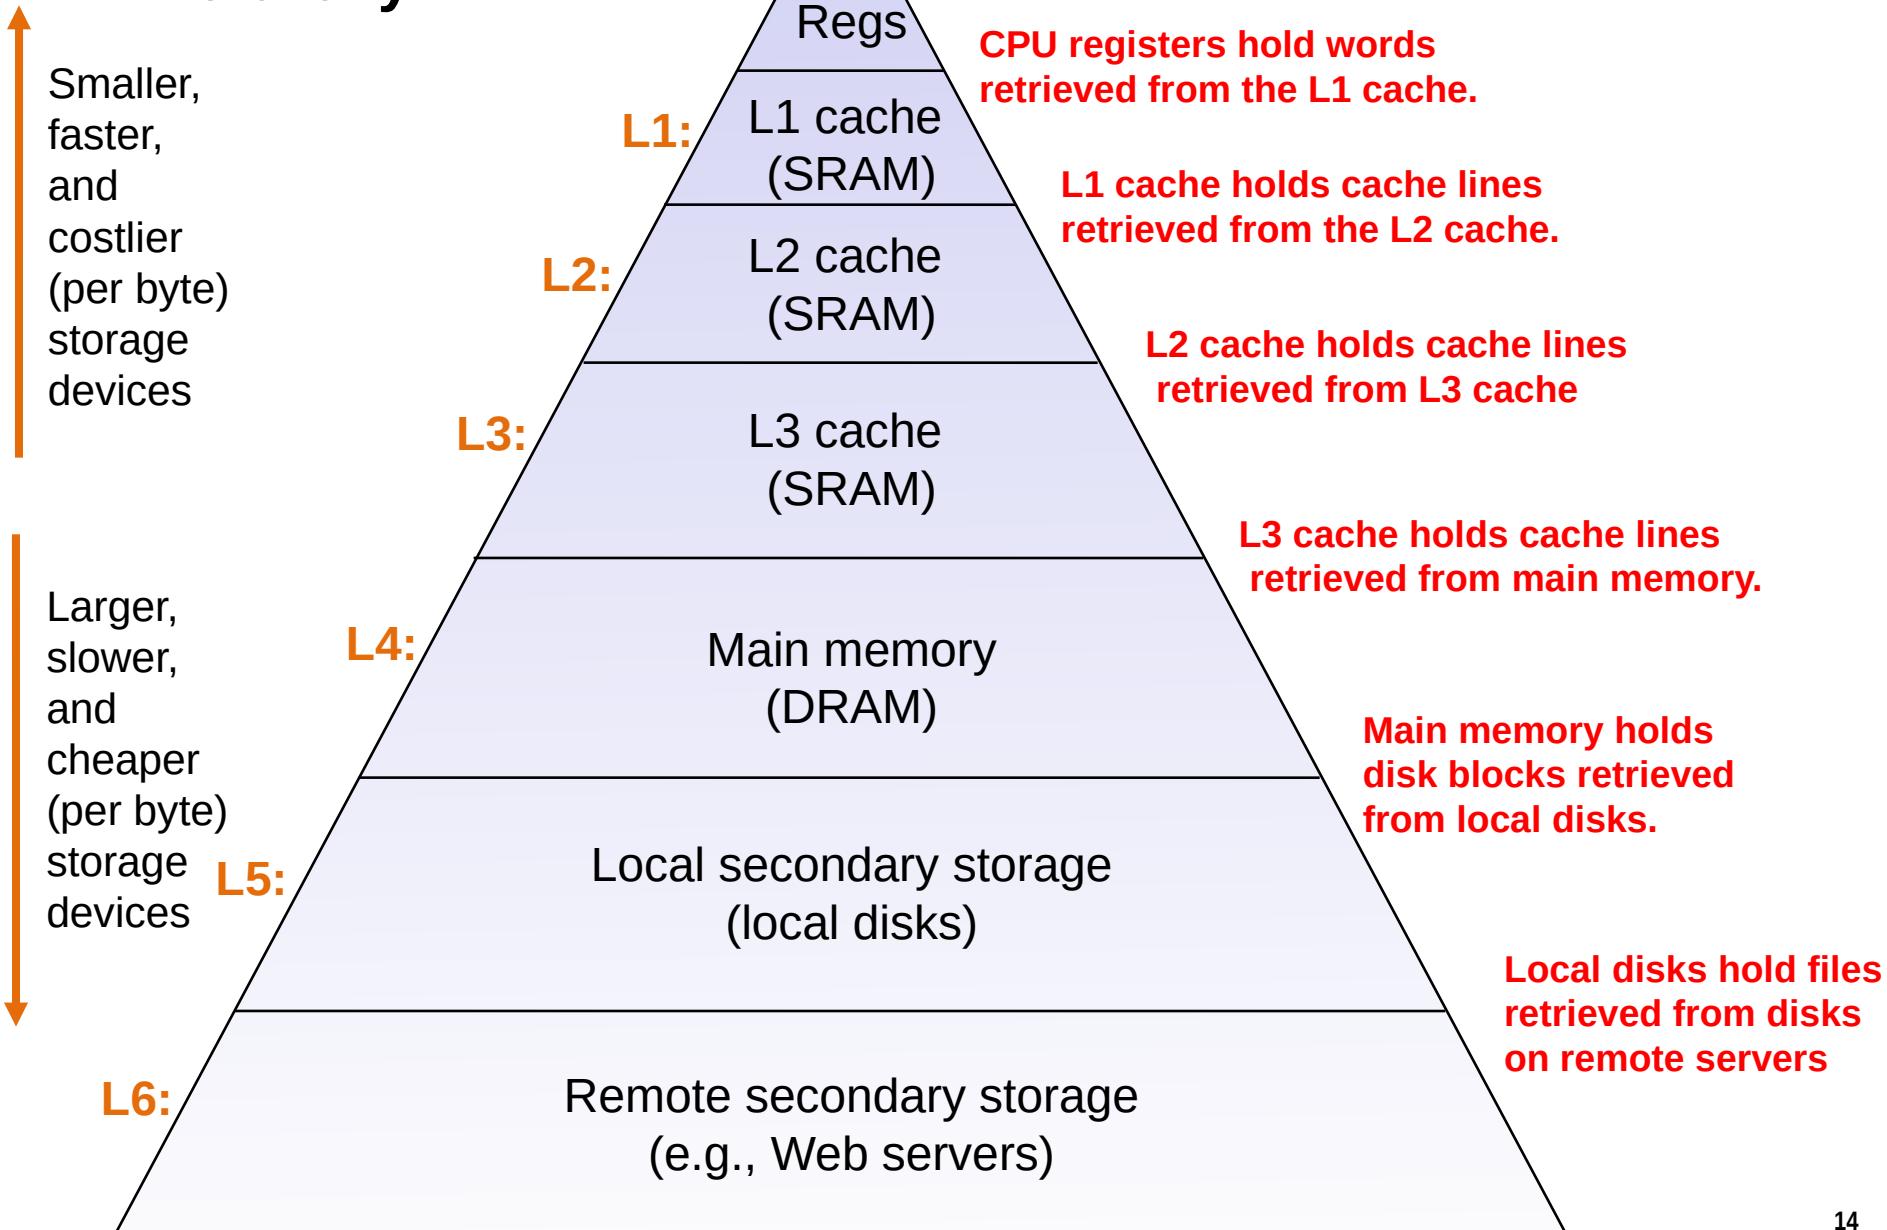
■ Instruction references

- Reference instructions in sequence.
- Cycle through loop repeatedly.

Spatial locality

Temporal locality

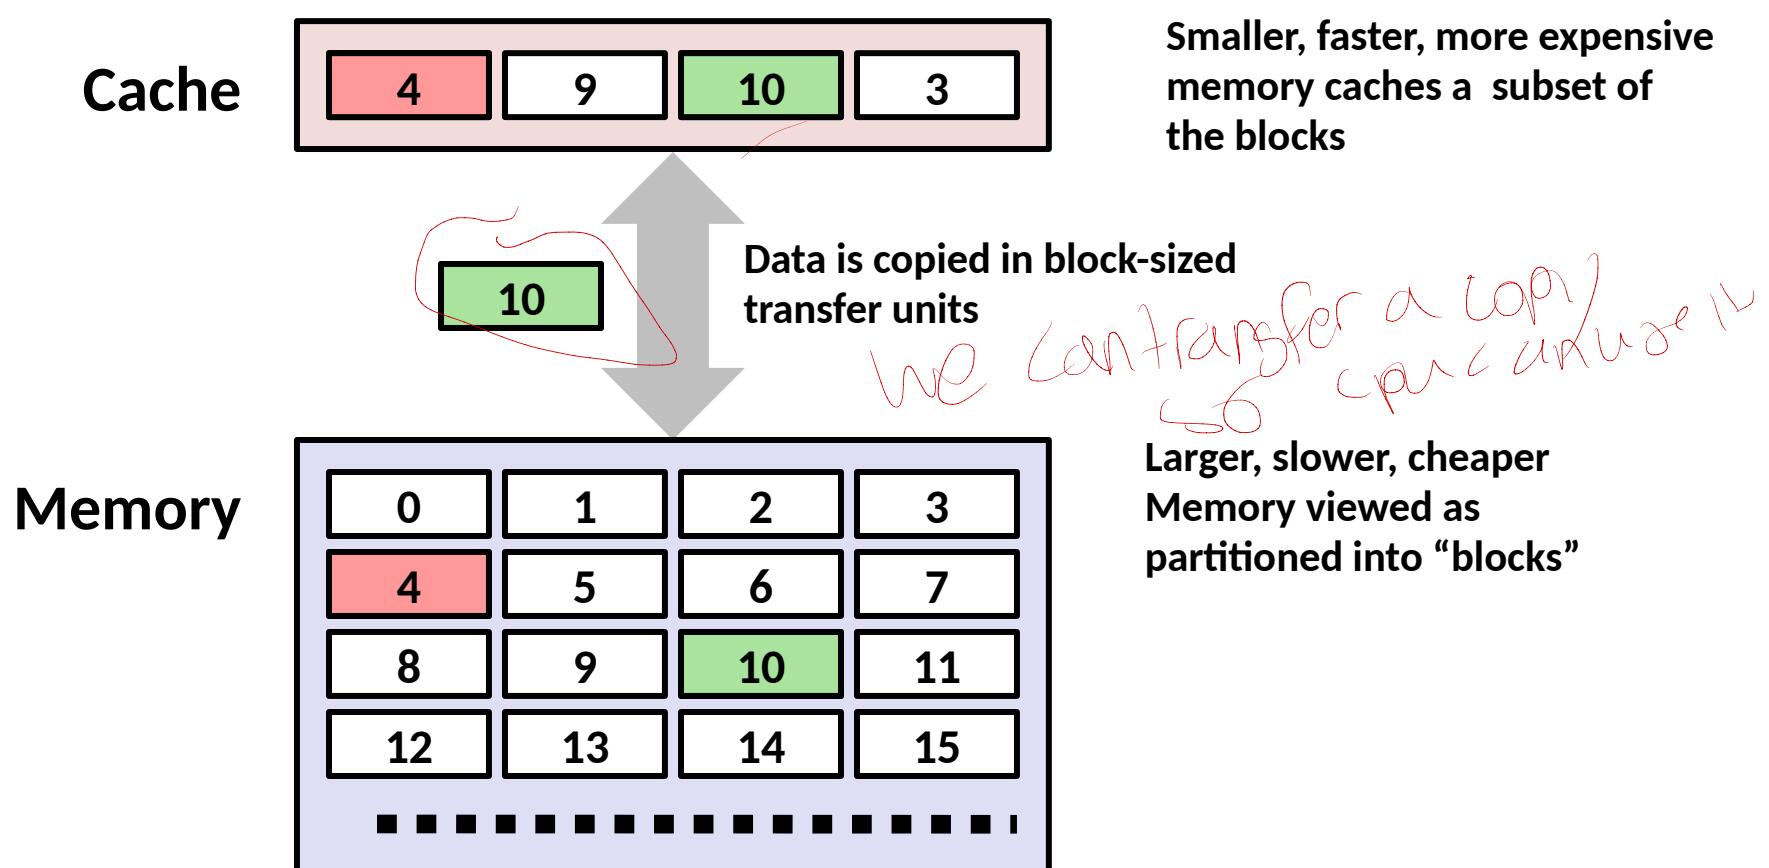
Example Memory Hierarchy



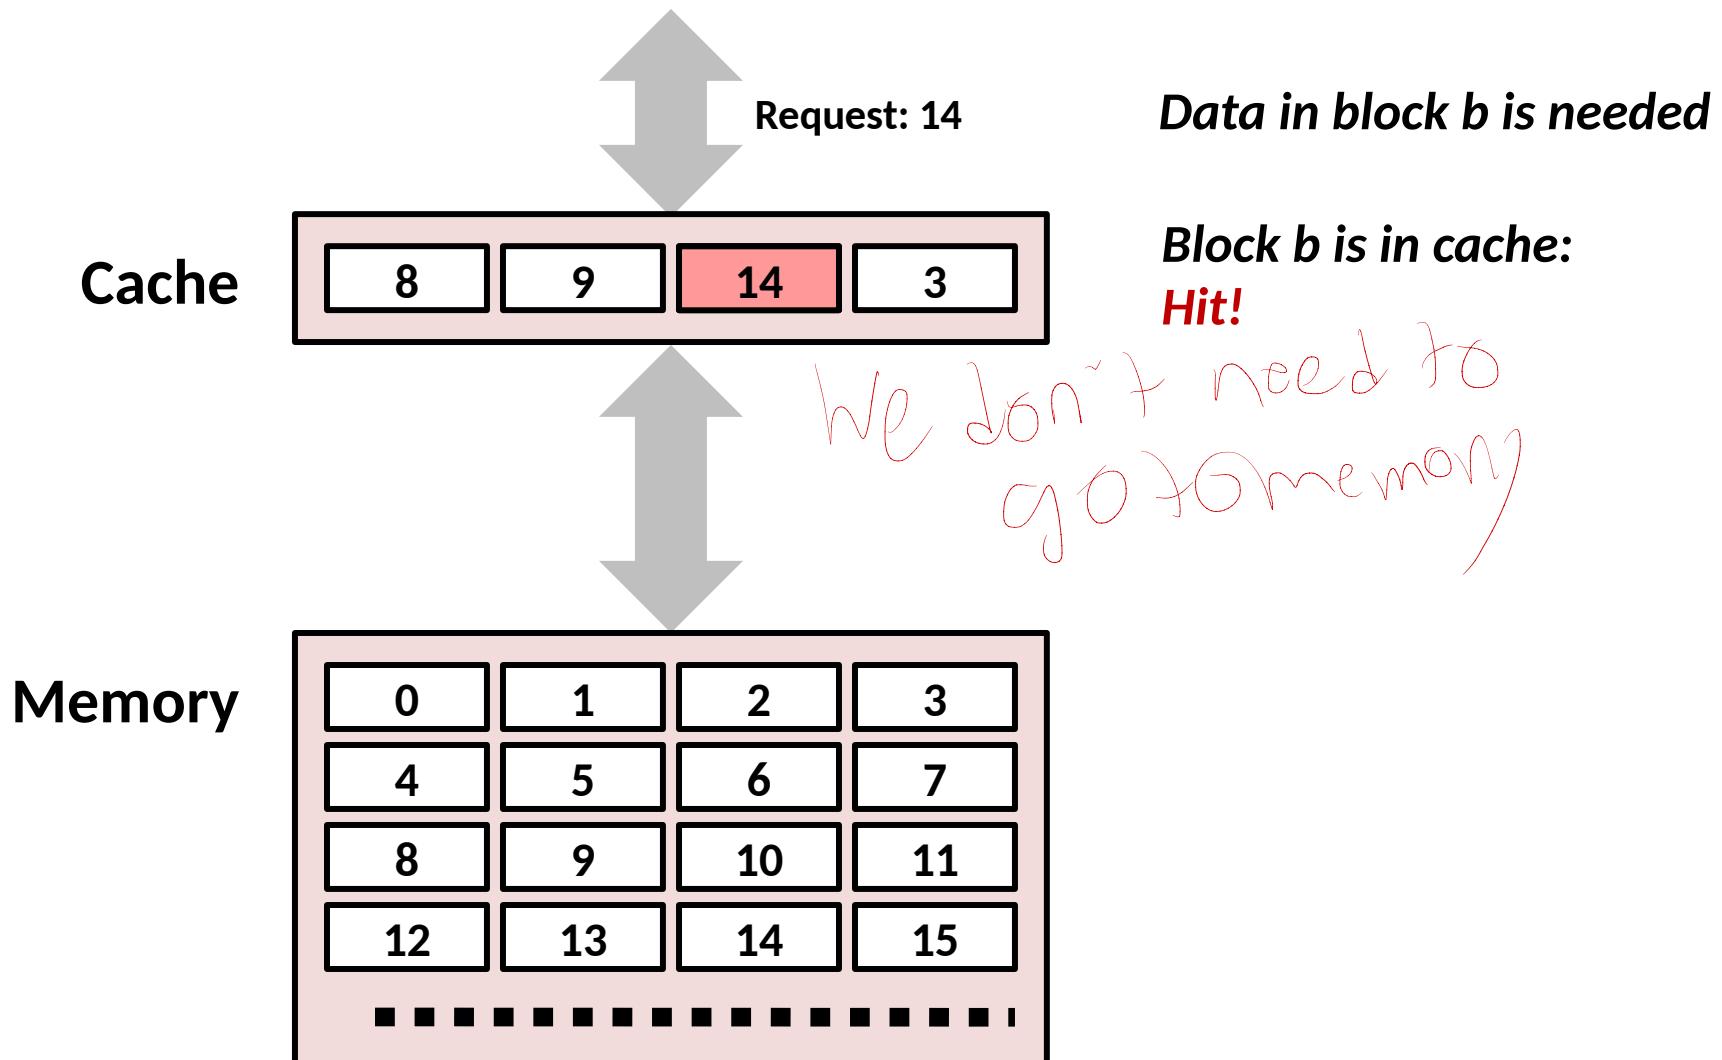
Caches

- ***Cache:*** A smaller, faster storage device that acts as a staging area for a subset of the data in a larger, slower device.
- Fundamental idea of a memory hierarchy:
 - For each k , the faster, smaller device at level k serves as a cache for the larger, slower device at level $k+1$.
- Why do memory hierarchies work?
 - Because of locality, programs tend to access the data at level k more often than they access the data at level $k+1$.
 - Thus, the storage at level $k+1$ can be slower, and thus larger and cheaper per bit.
- ***Big Idea:*** The memory hierarchy creates a large pool of storage that costs as much as the cheap storage near the bottom, but that serves data to programs at the rate of the fast storage near the top.

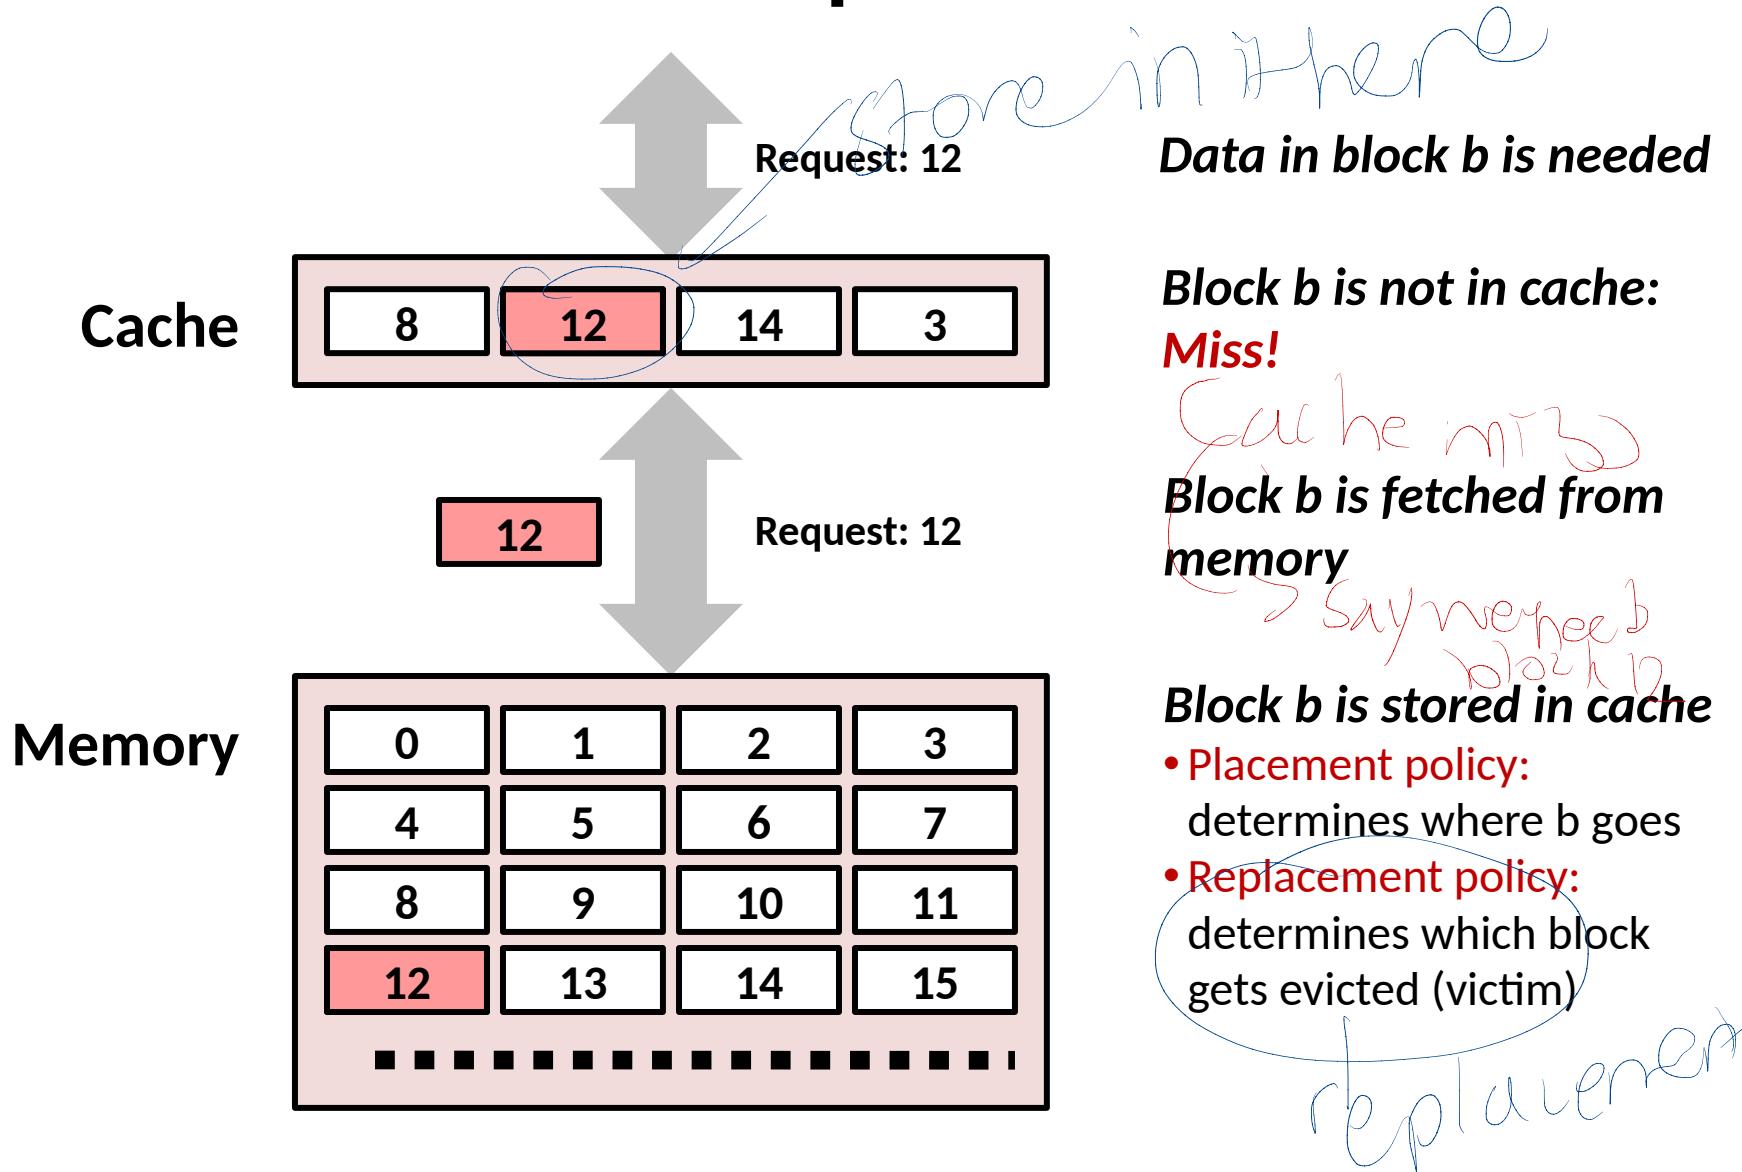
General Cache Concepts



General Cache Concepts: Hit



General Cache Concepts: Miss



General Caching Concepts:

Types of Cache Misses

■ Cold (compulsory) miss

- Cold misses occur because the cache is empty.

Charm it (get data)

■ Conflict miss

- Most caches limit blocks at level $k+1$ to a small subset (sometimes a singleton) of the block positions at level k .
 - E.g. Block i at level $k+1$ must be placed in block $(i \bmod 4)$ at level k .
 - Conflict misses occur when the level k cache is large enough, but multiple data objects all map to the same level k block.
 - E.g. Referencing blocks $0, 8, 0, 8, 0, 8, \dots$ would miss every time.

■ Capacity miss

- Occurs when the set of active cache blocks (**working set**) is larger than the cache.

↓
exit spots

require access to 10 memory

↑
put put
0 1 2 3
kick

Now up to store it

10 memory

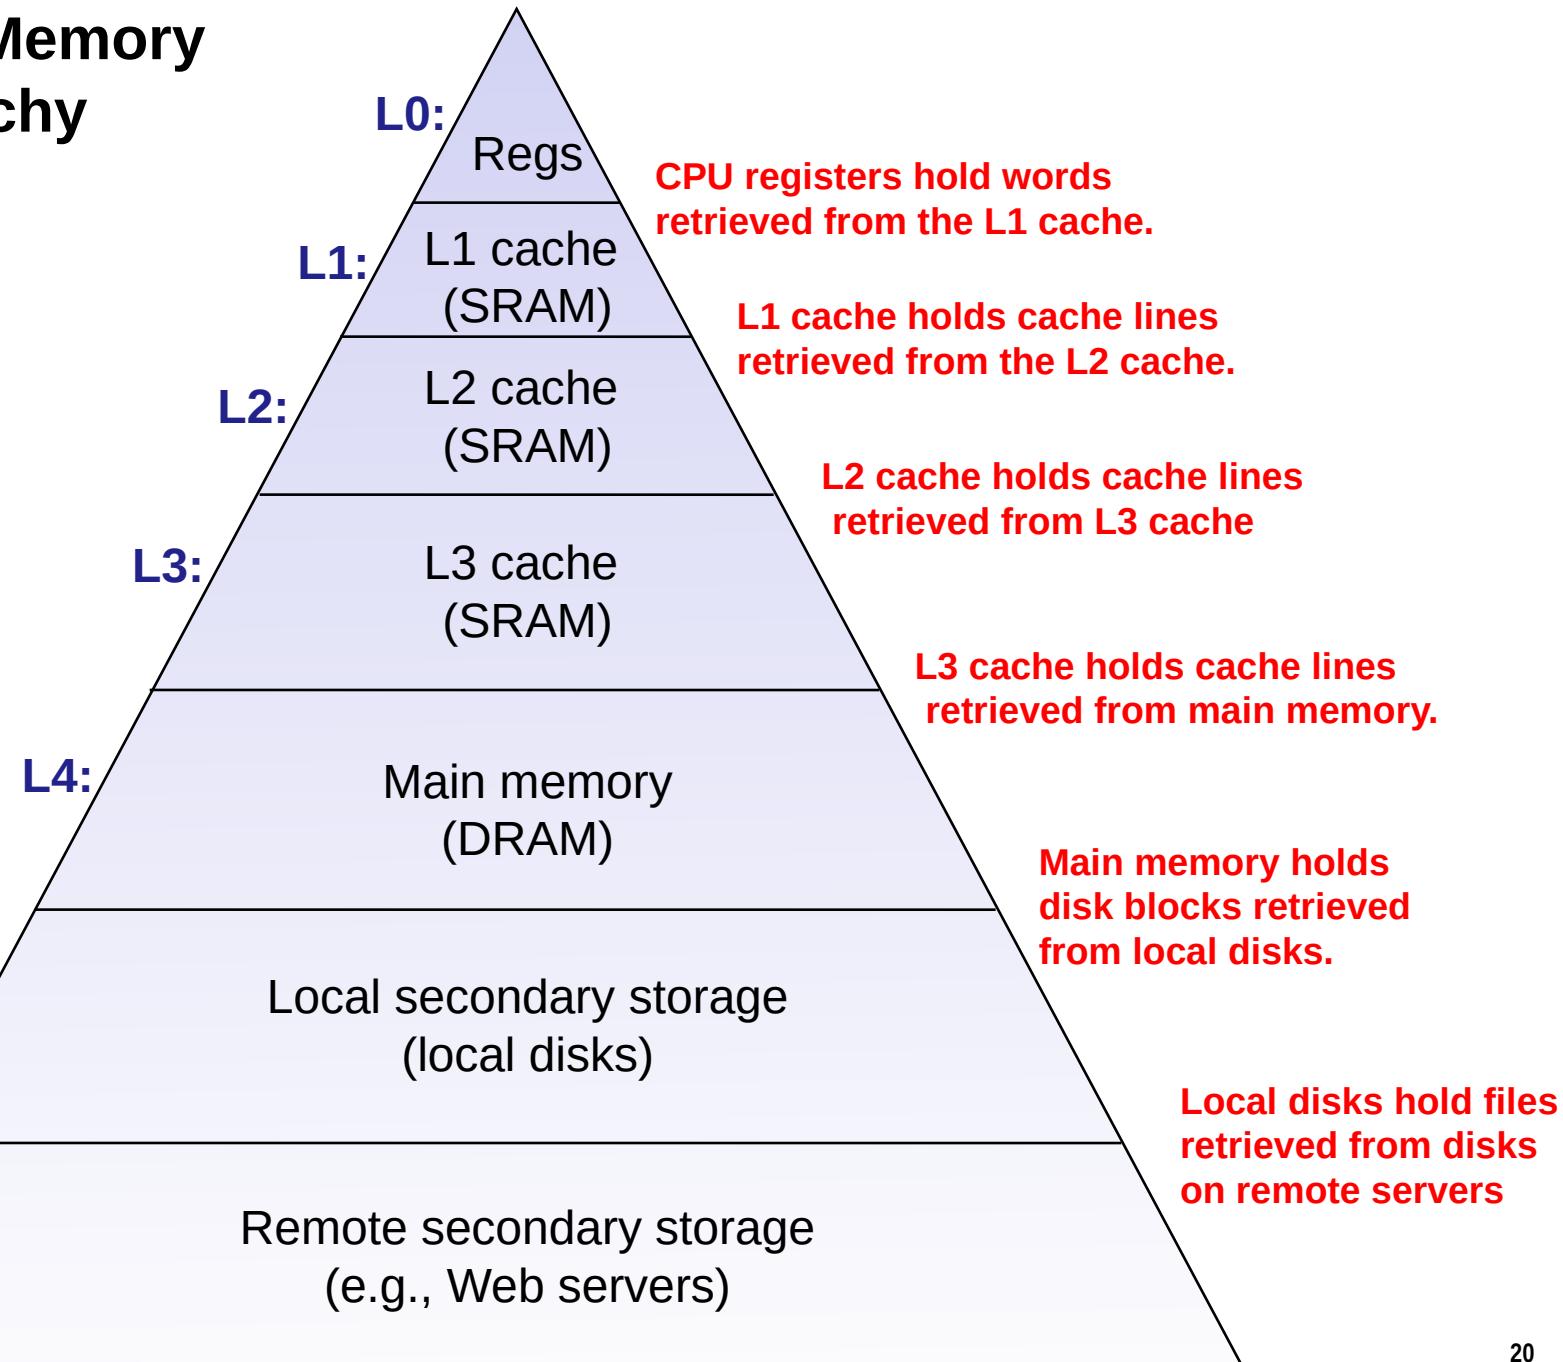
Example Memory Hierarchy

Smaller,
faster,
and
costlier
(per byte)
storage
devices

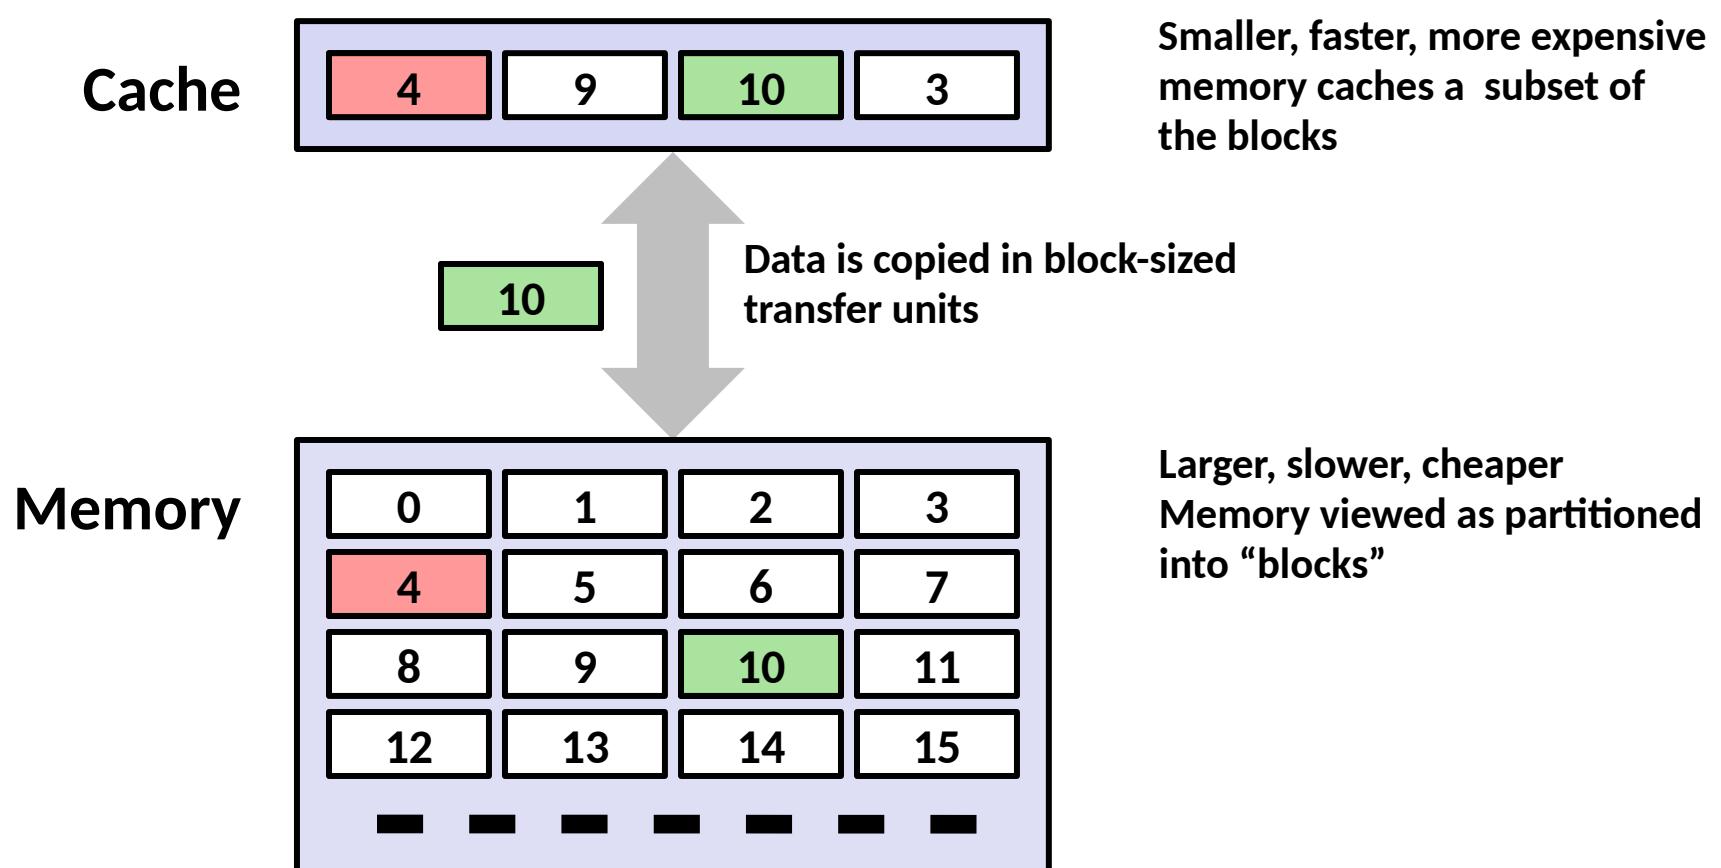
Larger,
slower,
and
cheaper
(per byte)
storage
devices

L6:

Remote secondary storage
(e.g., Web servers)

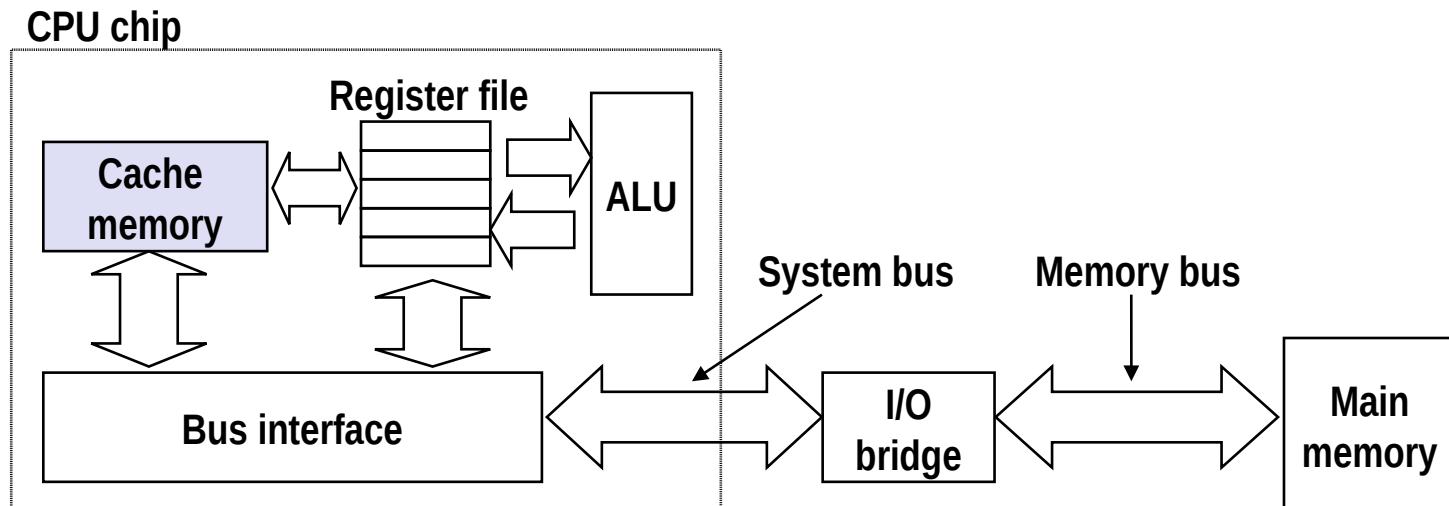


General Cache Concept

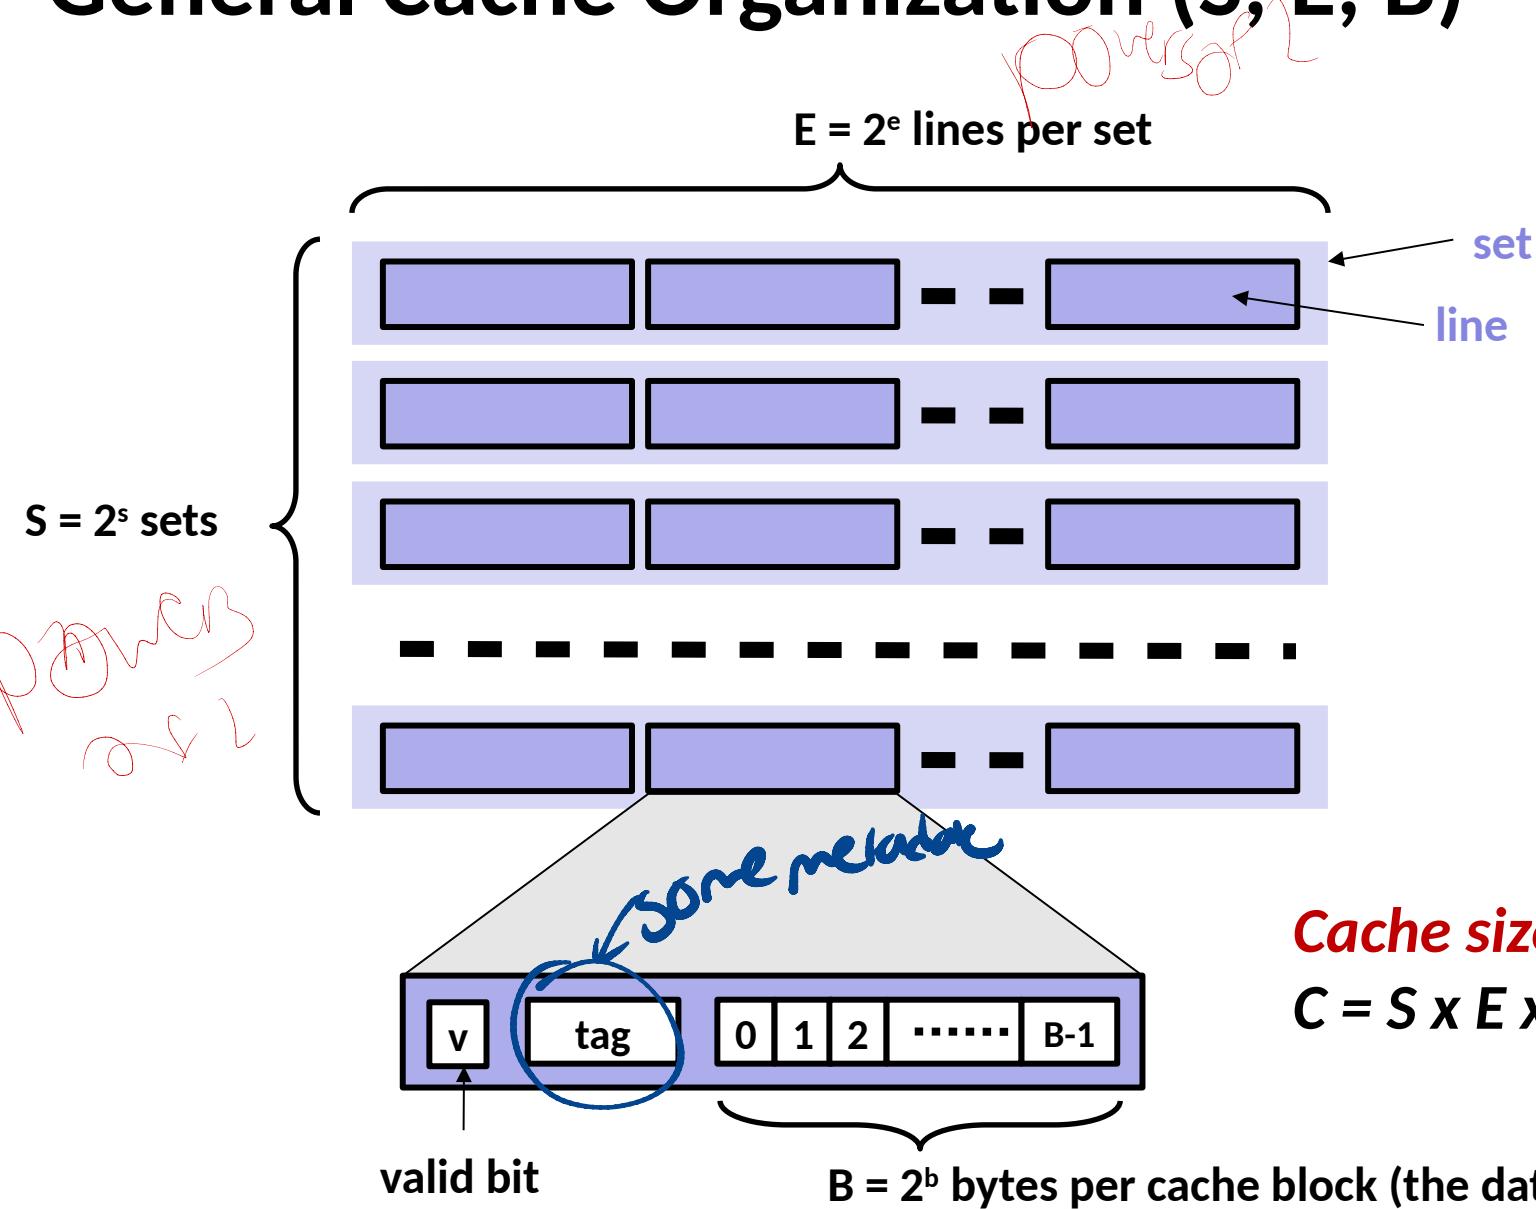


Cache Memories

- Cache memories are small, fast SRAM-based memories managed automatically in hardware
 - Hold frequently accessed blocks of main memory
- CPU looks first for data in cache
- Typical system structure:



General Cache Organization (S, E, B)

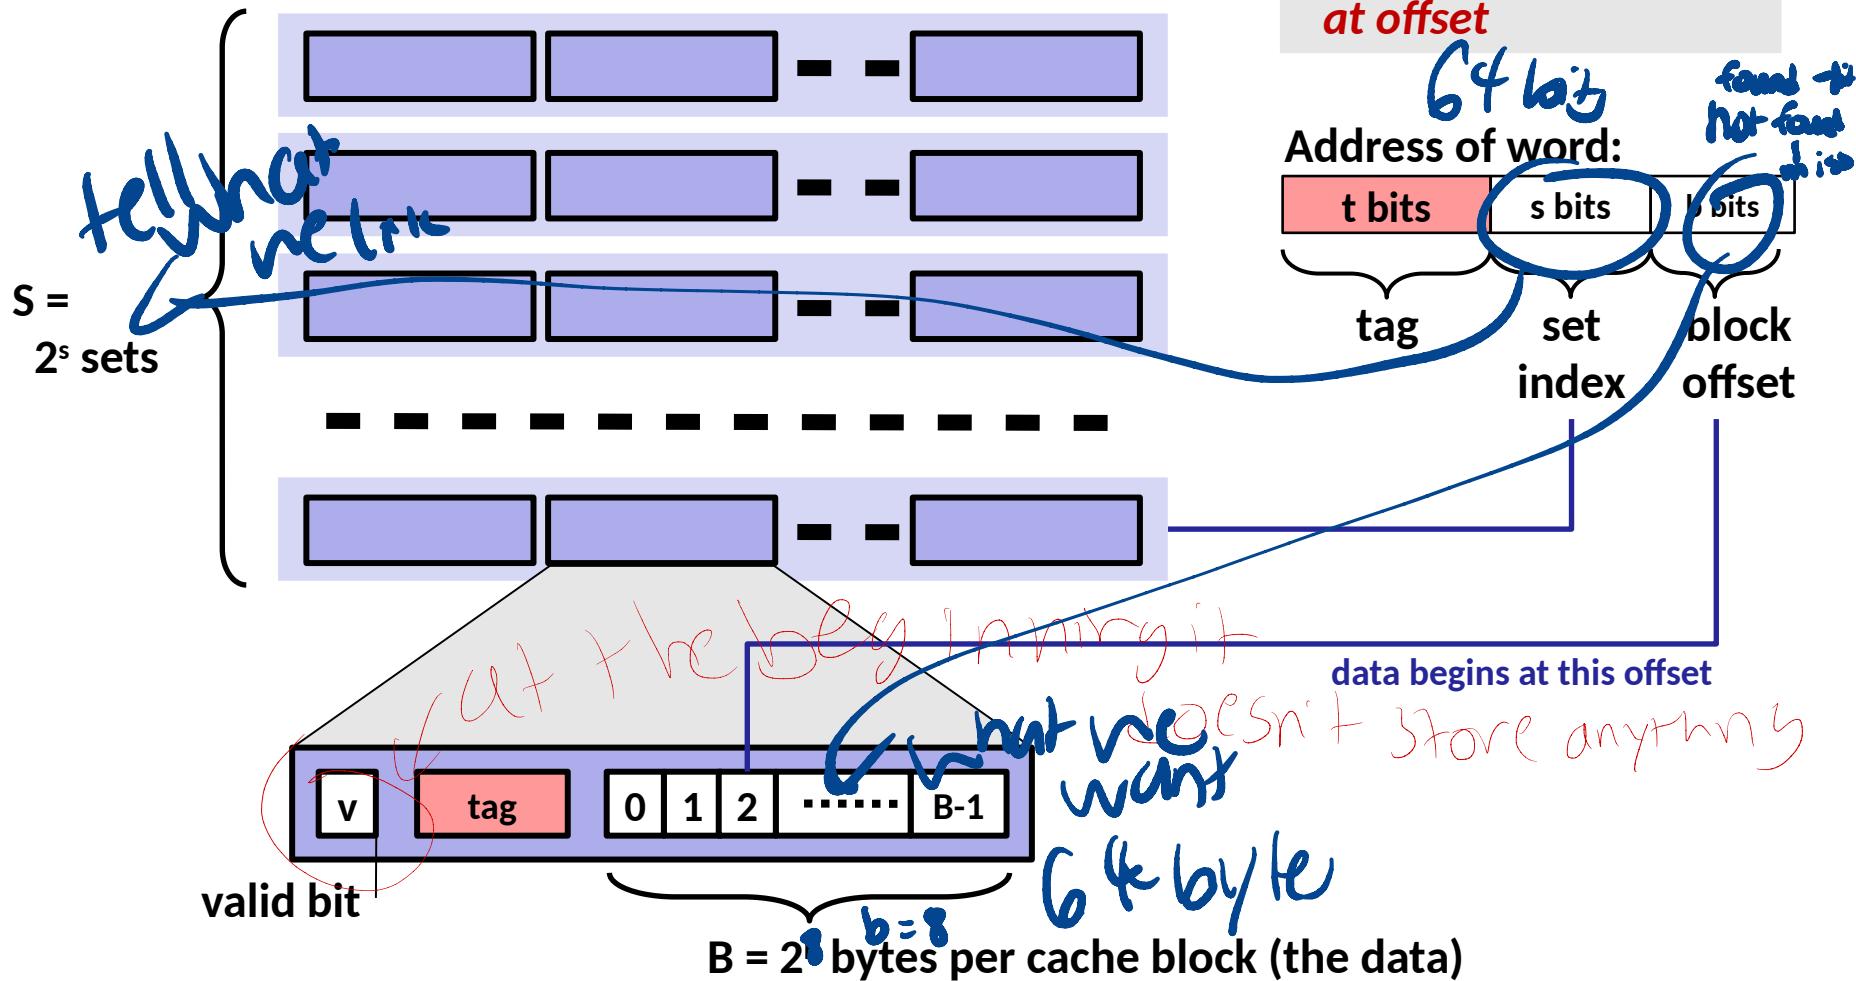


Cache Read

Steps

1. s bits - match is a hit
 2. tag - determine what to return
 3. block offset
- $E = 2^e$ lines per set

- Locate set
- Check if any line in set has matching tag
- Yes + line valid: hit
- Locate data starting at offset

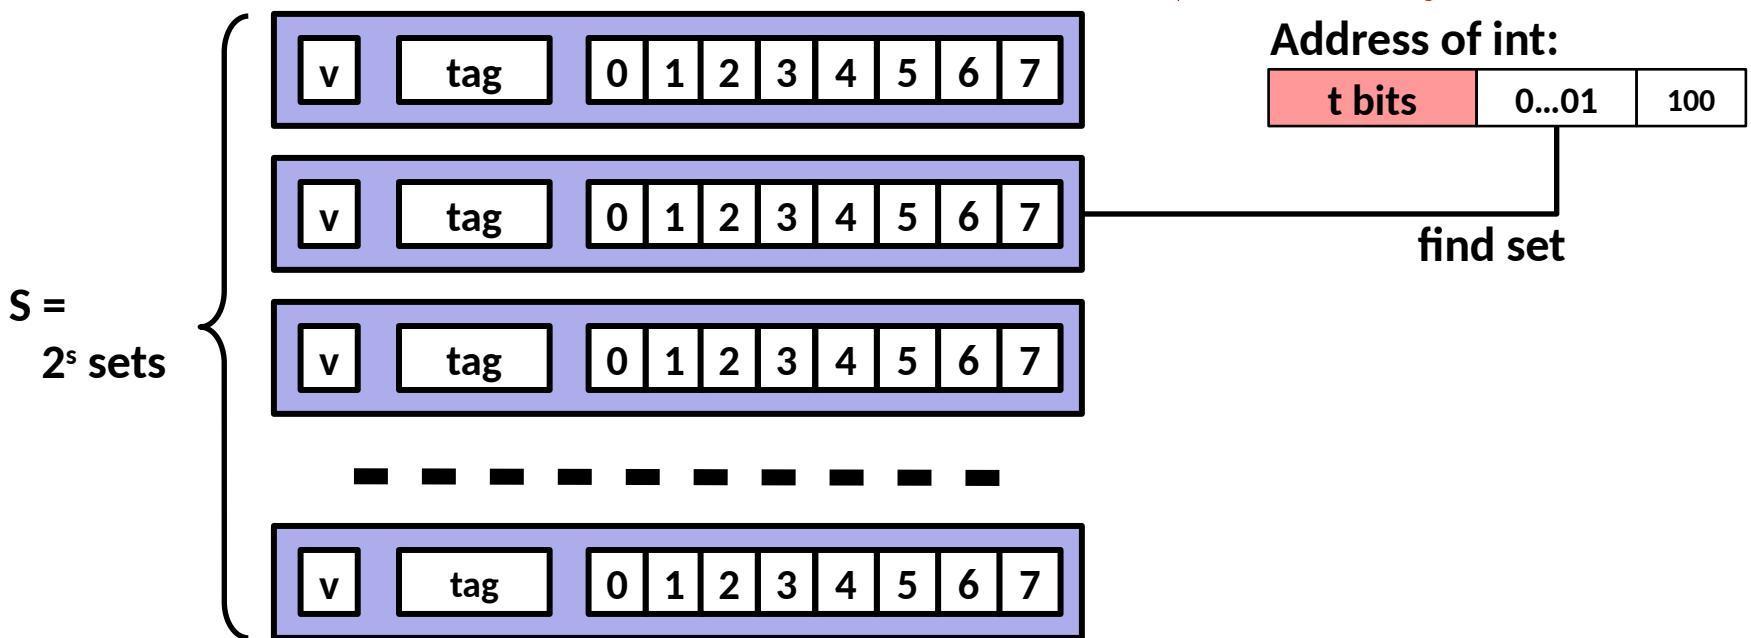


Ex.: Direct Mapped Cache ($E = 1$)

Direct mapped: One line per set

Assume: cache block size 8 bytes

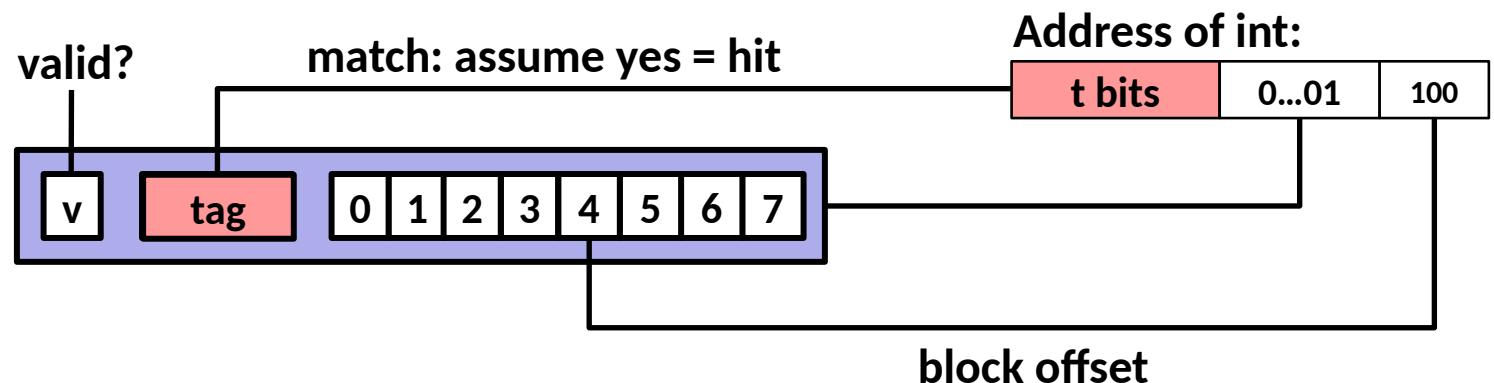
add a single line
in each set



Ex.: Direct Mapped Cache ($E = 1$)

Direct mapped: One line per set

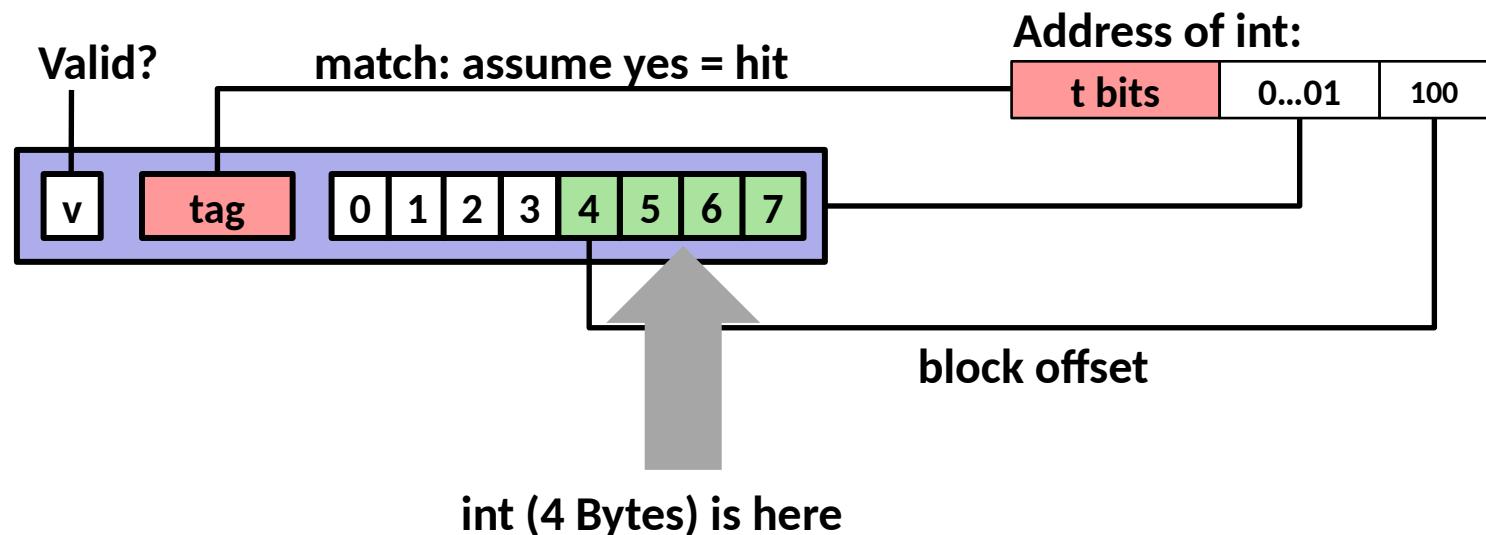
Assume: cache block size 8 bytes



Ex.: Direct Mapped Cache ($E = 1$)

Direct mapped: One line per set

Assume: cache block size 8 bytes



If tag doesn't match:
old line is evicted and replaced

Direct-Mapped Cache Simulation

$t=1$	$s=2$	$b=1$
x	xx	x

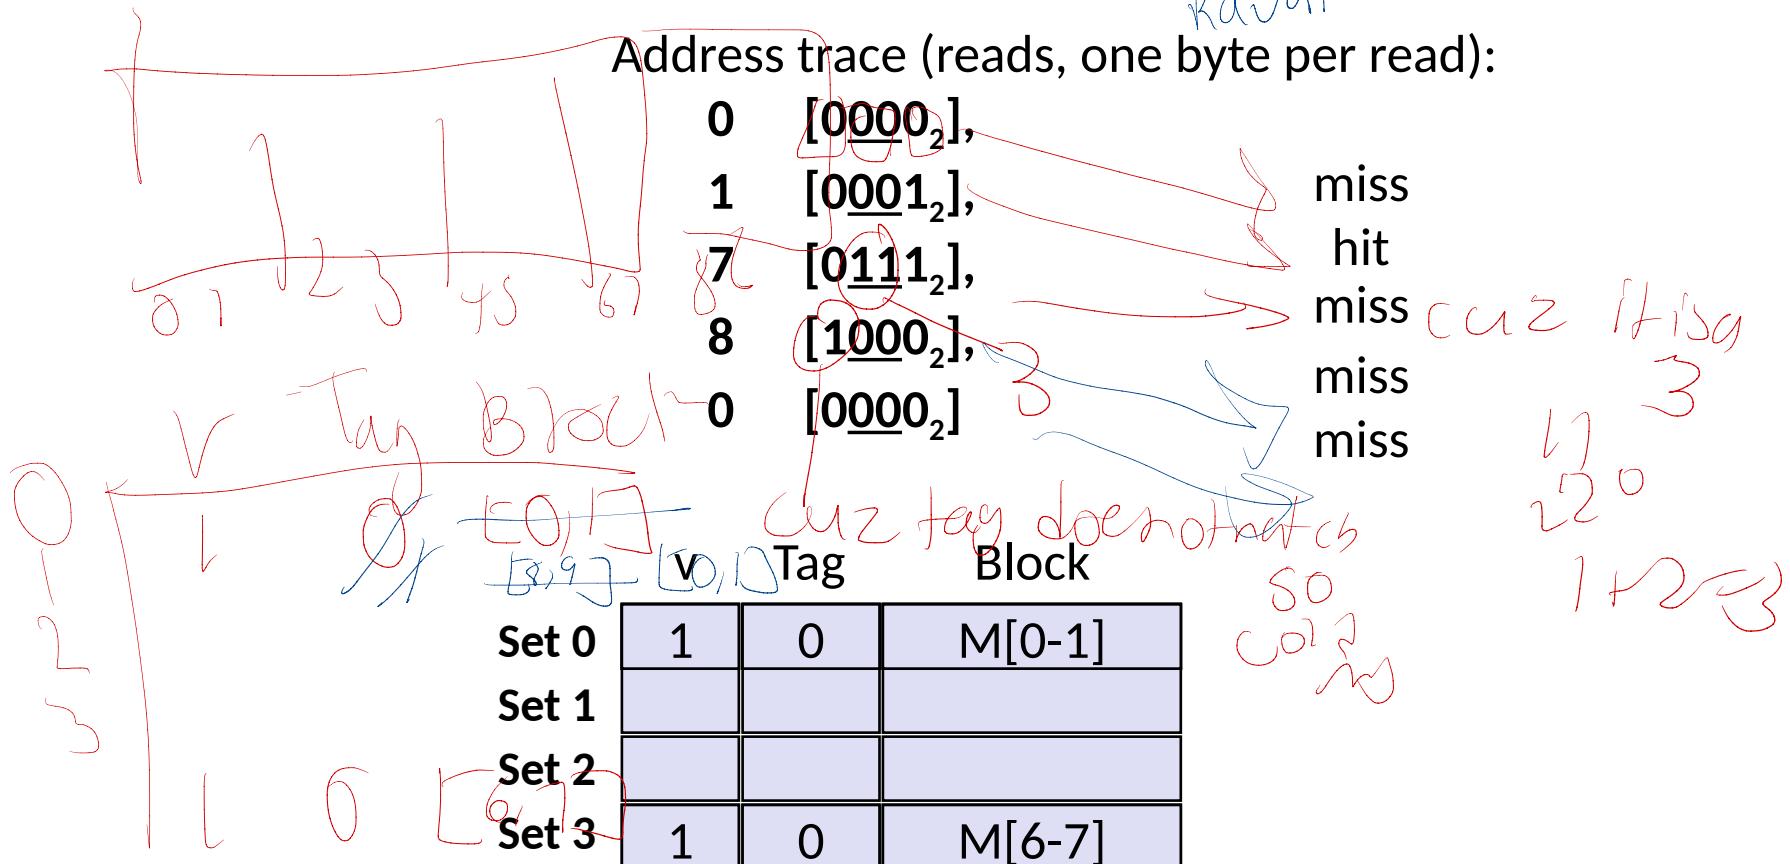
$M=16$ bytes (4-bit addresses), $B=2$ bytes/block,
 $S=4$ sets, $E=1$ Blocks/set

Siraj
Raval

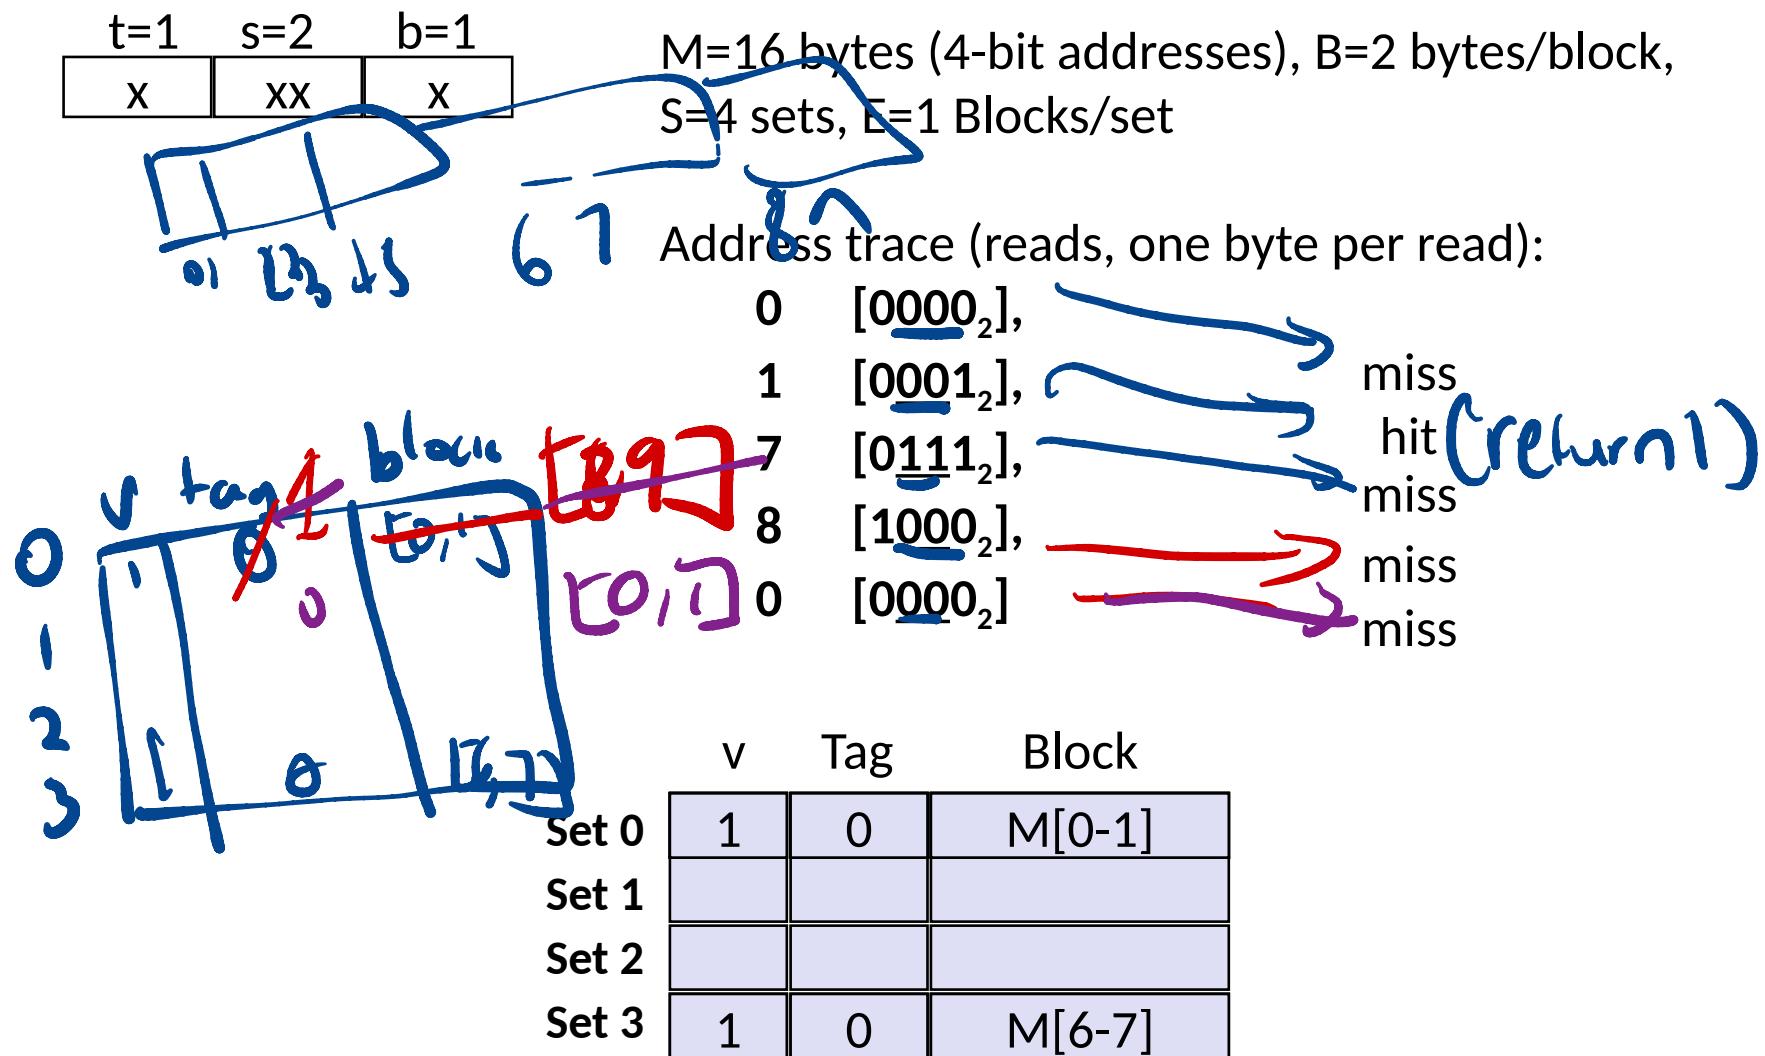
Address trace (reads, one byte per read):

- 0 [0000]₂,
- 1 [0001]₂,
- 7 [0111]₂,
- 8 [1000]₂,
- 0 [0000]₂

miss
hit
miss
miss
miss



Direct-Mapped Cache Simulation

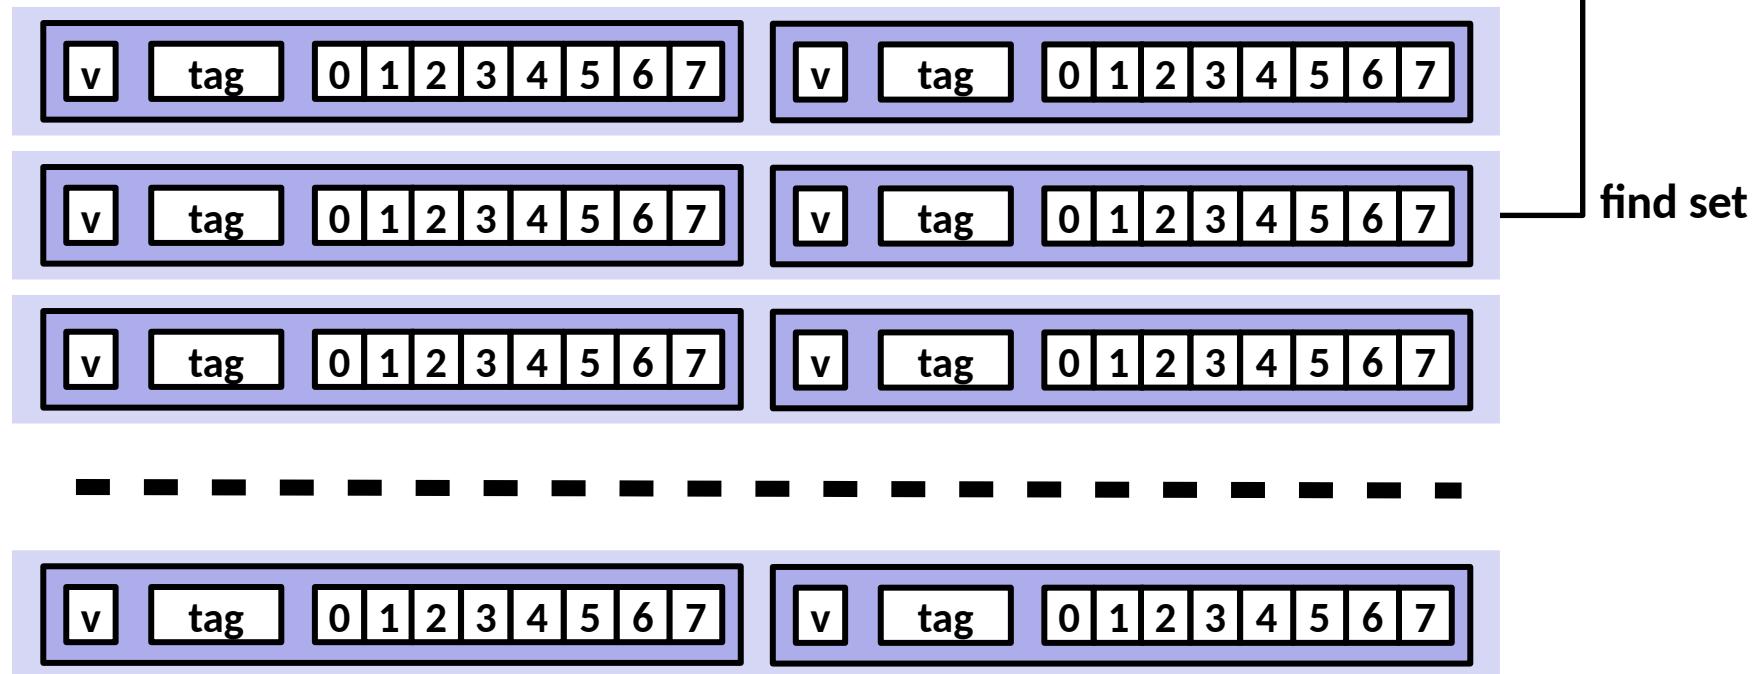
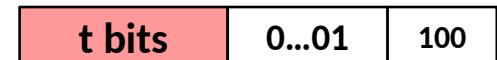


E-way Set Associative Cache (Here: E = 2)

E = 2: Two lines per set

Assume: cache block size 8 bytes

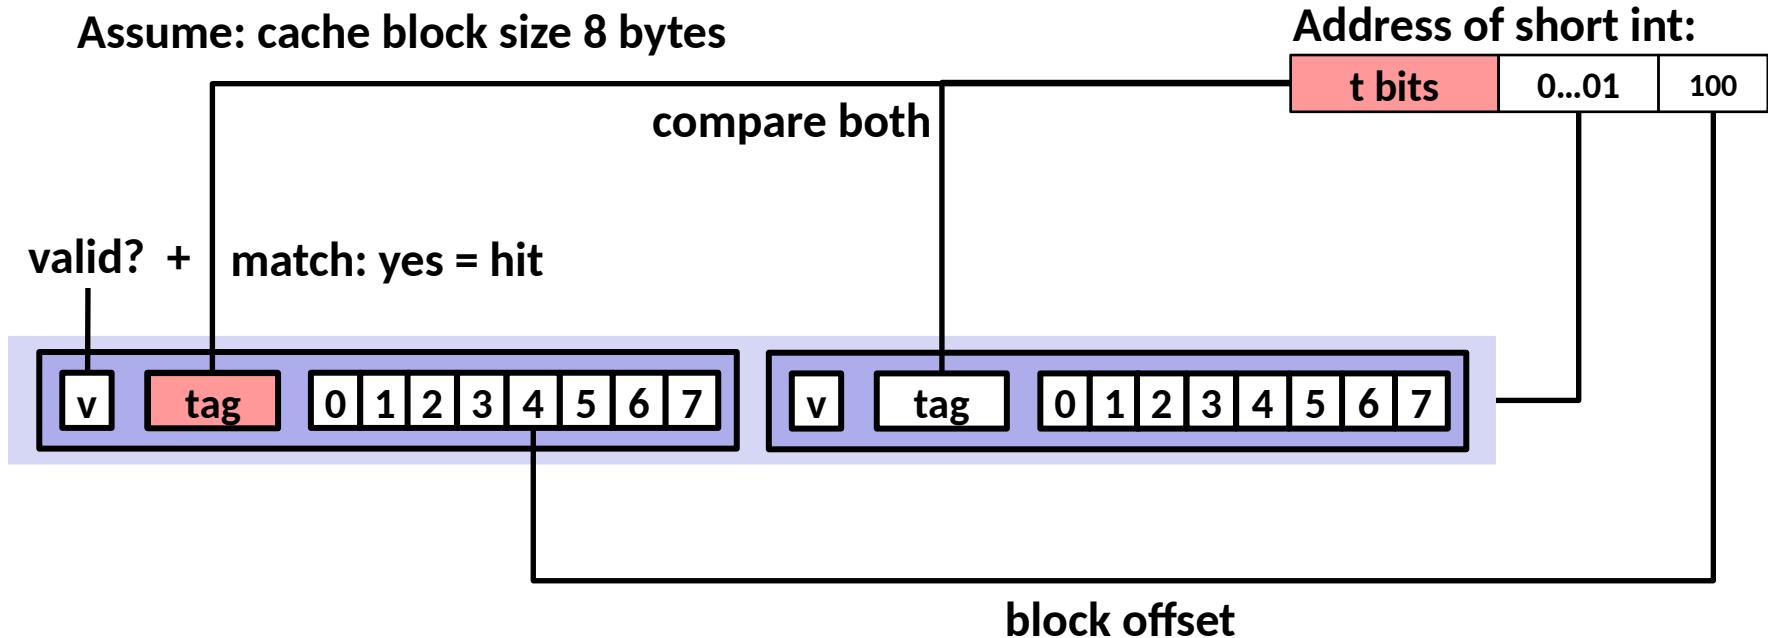
Address of short int:



E-way Set Associative Cache (Here: E = 2)

E = 2: Two lines per set

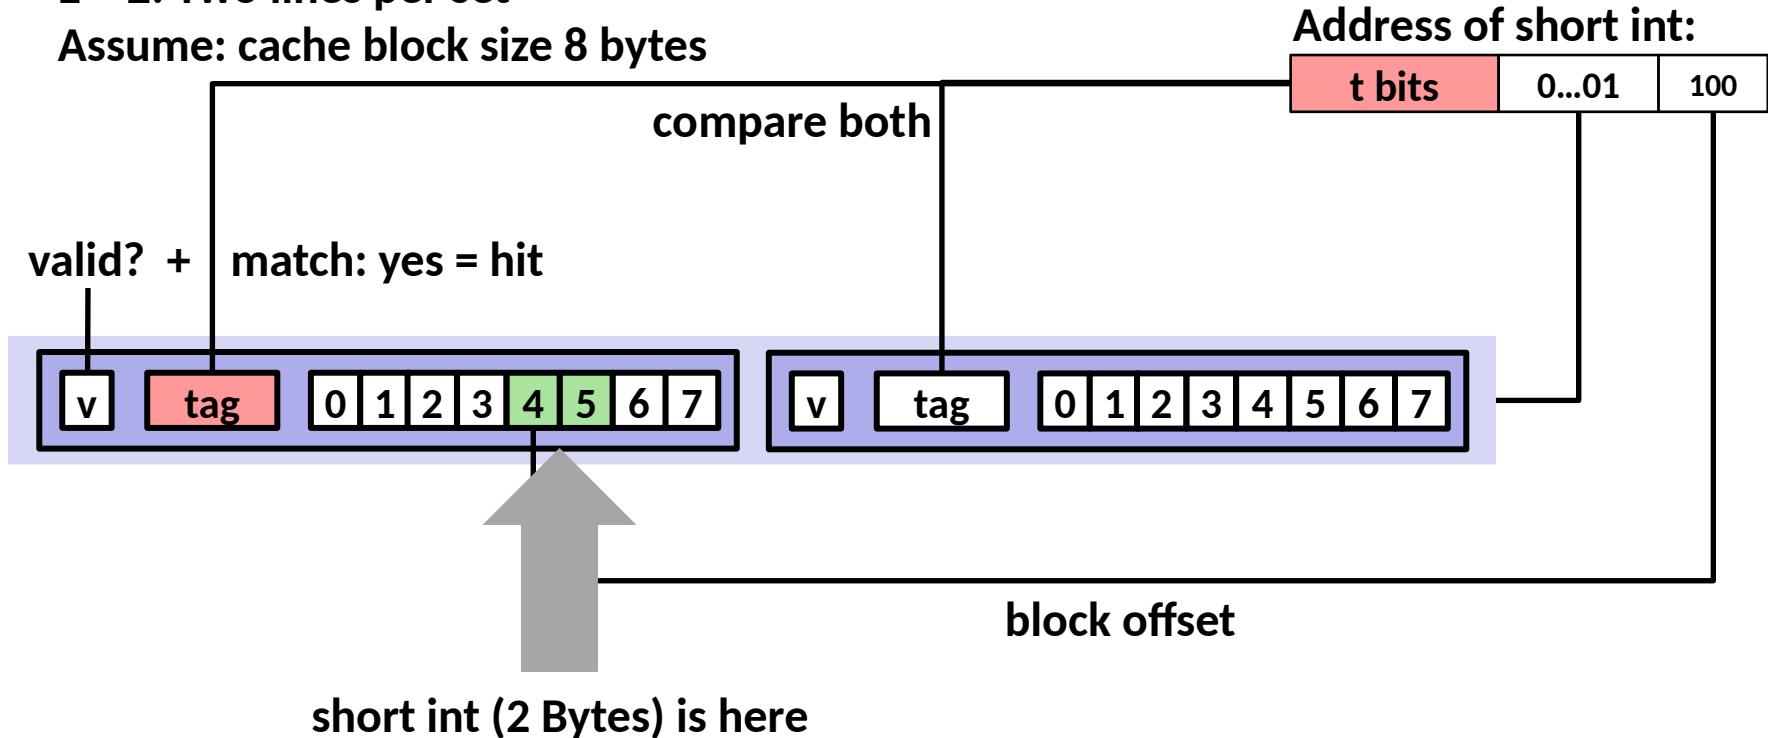
Assume: cache block size 8 bytes



E-way Set Associative Cache (Here: E = 2)

E = 2: Two lines per set

Assume: cache block size 8 bytes



No match:

- One line in set is selected for eviction and replacement
- Replacement policies: random, least recently used (LRU), ...

2-Way Set Associative Cache Simulation

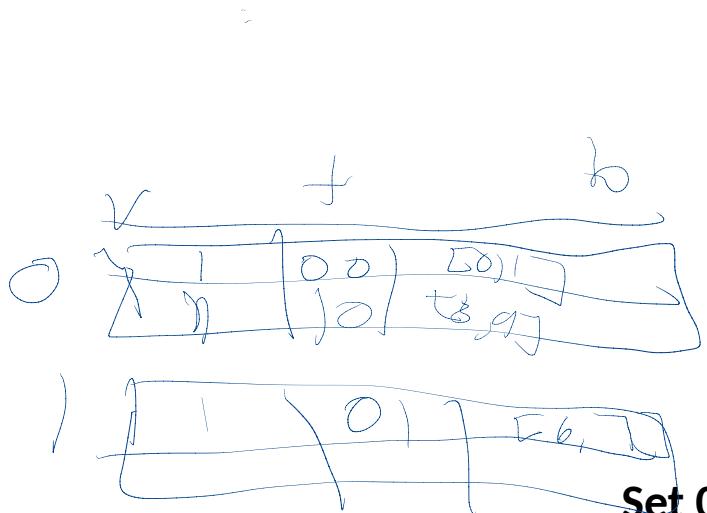
$t=2 \quad s=1 \quad b=1$

xx	x	x
----	---	---

$M=16$ byte addresses, $B=2$ bytes/block,
 $S=2$ sets, $E=2$ blocks/set

Address trace (reads, one byte per read):

0	[0000 ₂],	miss
1	[0001 ₂],	hit
7	[0111 ₂],	miss
8	[1000 ₂],	miss
0	[0000 ₂]	hit



	v	Tag	Block
Set 0	1	00	M[0-1]
	1	10	M[8-9]
Set 1	1	01	M[6-7]
	0		

2-Way Set Associative Cache Simulation

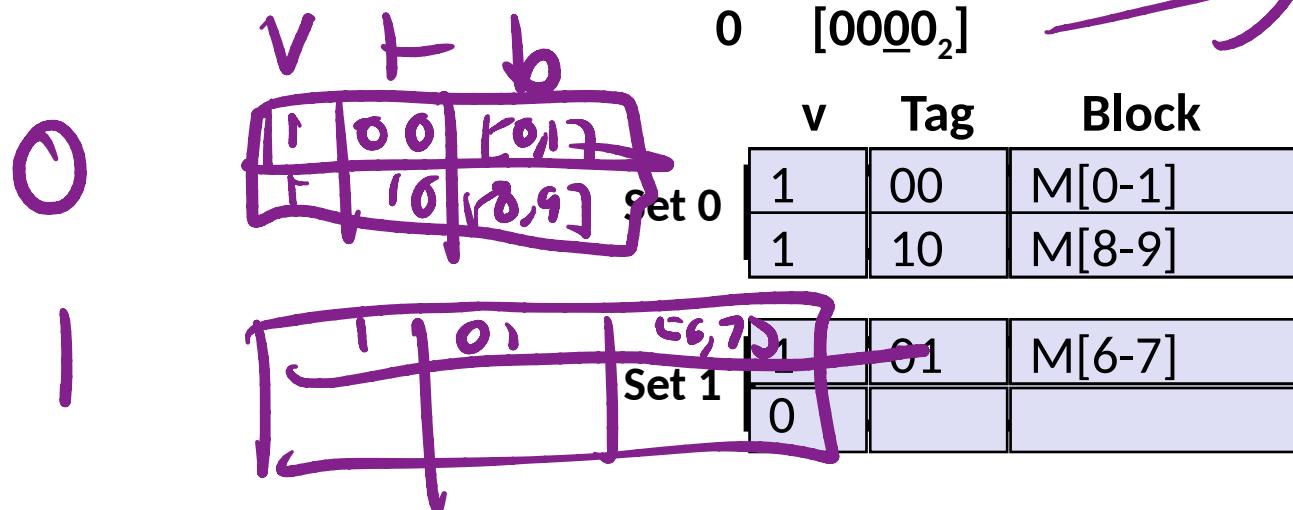
$t=2$ $s=1$ $b=1$

xx	x	x
----	---	---

$M=16$ byte addresses, $B=2$ bytes/block,
 $S=2$ sets, $E=2$ blocks/set

Address trace (reads, one byte per read):

0	[0000 ₂],	miss
1	[0001 ₂],	hit
7	[0111 ₂],	miss
8	[1000 ₂],	miss
0	[0000 ₂]	hit

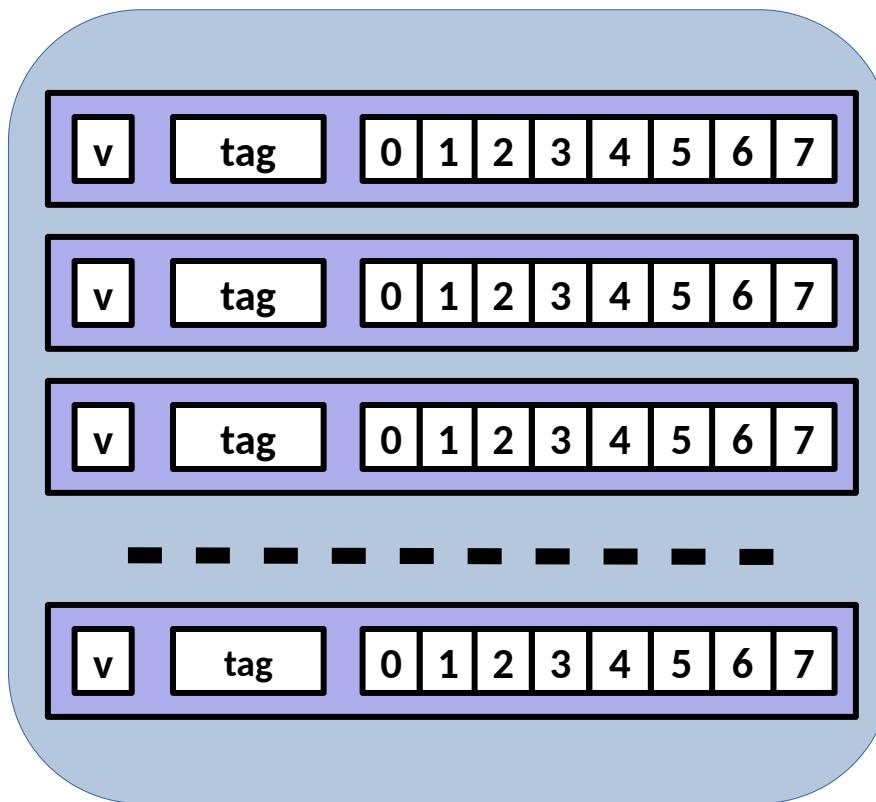


Ex.: Fully Associative Cache ($S = 1$)

One set containing all lines

Assume: cache block size 8 bytes

One set:



Address of int:



Cache associativity

- Direct mapped
 - $E = 1$ (1 line per set)
- n-way associative
 - n lines per set
- Fully associative
 - $S = 1$ (1 set containing all lines)

Cache associativity

- Requested address gives all information needed to locate data in the cache:

- Set index: which bookshelf?

↙ \Rightarrow Set

- Tag: book title

↙ \Rightarrow Tag

- Block index: what page of the book?

↙ \Rightarrow Block

Cache Size

- $C = S * E * B$
 - S = #sets
 - E = #lines / set (associativity)
 - B = #bytes / block (in a line)
- Suppose we're given:
 - 32 KB cache
 - 64 B line
 - 4-way set associative
 - 32-bit architecture
- How many tag, set, block index bits?

Cache Size

- Suppose we're given:
 - 32 KB cache
 - 64 B line
 - 4-way set associative
 - 32-bit architecture
- How many tag, set, block index bits?
 - 64 B per block = ? bits
 - How many lines in a 32 KB cache?
 - How many sets to hold those lines?

$$\# \text{ Sets} = \frac{32 \text{ KB}}{64 \text{ B}} = 128 = 2^7$$

$$2^6 = 64$$
$$\frac{32 \text{ KB}}{64 \text{ B}} = \frac{2^5 \cdot 2^{10}}{2^6} = 2^9 = ST$$

Cache Size

■ Suppose we're given:

- 32 KB cache
- 64 B line
- 4-way set associative
- 32-bit architecture

13
2

■ How many tag, set, block index bits?

- 64 B per block = 6 bits
- 32 KB cache / 64 B per line = 2^9 lines
- 2^9 lines / 4 lines per set = 2^7 sets
- So 6 block index bits, 7 set index bits, ??? tag bits



$$19 + 7 + 6 = 32$$

64 B 2^9 LNS

Cache Size

■ Suppose we're given:

- 32 KB cache
- 64 B line
- 4-way set associative
- 32-bit architecture

■ How many tag, set, block index bits?

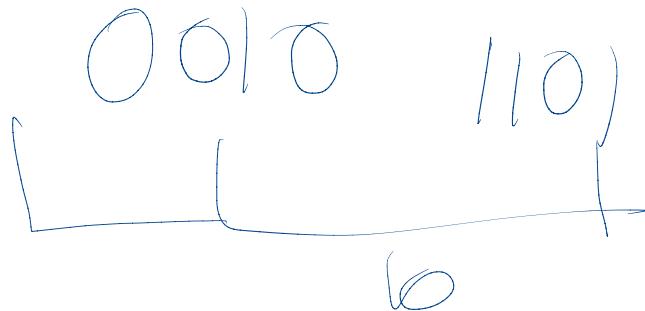
- 64 B per block = 6 bits
- 32 KB cache / 64 B per line = 2^9 lines
- 2^9 lines / 4 lines per set = 2^7 sets
- So 6 block index bits, 7 set index bits
- $32 - 6 - 7 = 19$ tag bits

32 line
0.75

Cache Size

- Suppose we're given:
 - Read 0x2b74ce2d *figure if it's in our cache*
- What are the tag, set, block index?
 - (19 tag, 7 set, 6 block)

0010 1011 0111 0100 1100 1110

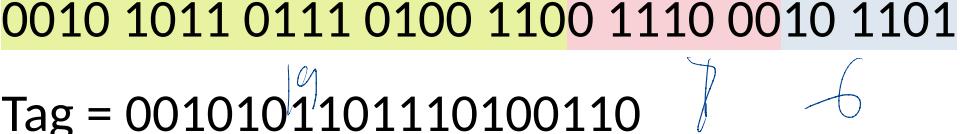


Cache Size

- Suppose we're given:
 - Read 0x2b74ce2d
- What are the tag, set, block index?
 - (19 tag, 7 set, 6 block)
 - 0x2b74ce2d =

0010 1011 0111 0100 1100 1110 0010 1101

Cache Size

- Suppose we're given:
 - Read 0x2b74ce2d
- What are the tag, set, block index?
 - (19 tag, 7 set, 6 block)
 - 0x2b74ce2d =


The binary string is 0010 1011 0111 0100 1100 1110 0010 1101. It is divided into three colored segments: yellow (bits 0-18), red (bits 19-25), and blue (bits 26-31). Handwritten indices are shown above the string: '19' above the yellow segment, '7' above the red segment, and '6' above the blue segment.

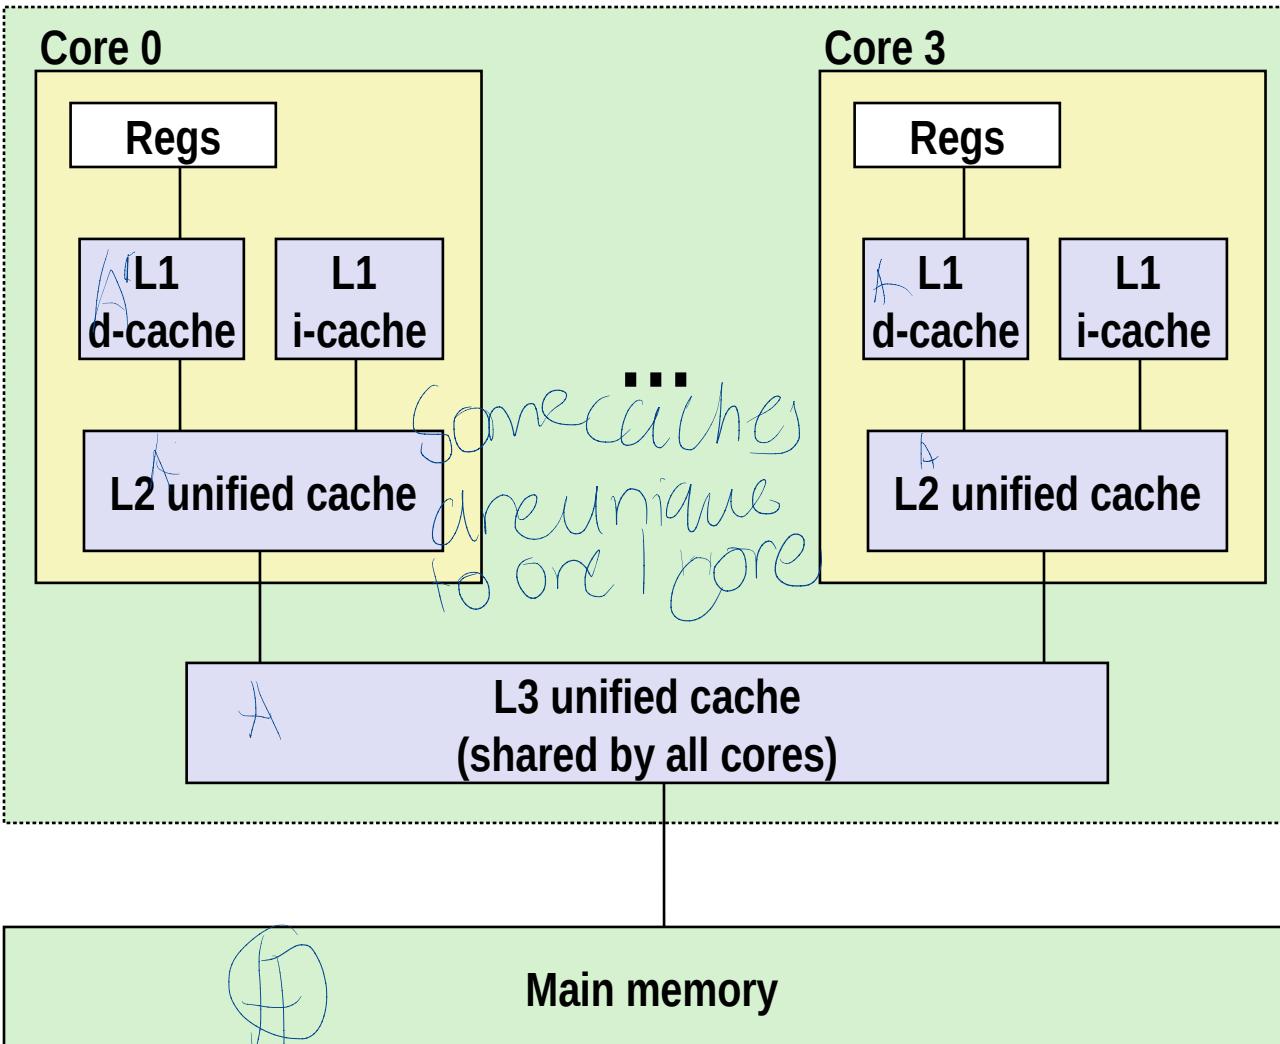
 - Tag = 0010101101110100110
 - Set index = 0111000 (56)
 - Block index = 101101 (45)

What about writes?

- **Multiple copies of data exist:**
 - L1, L2, L3, Main Memory, Disk
- **What to do on a write-hit?**
 - Write-through (write immediately to memory)
 - Write-back (defer write to memory until replacement of line)
 - Need a dirty bit (line different from memory or not)
- **What to do on a write-miss?**
 - Write-allocate (load into cache, update line in cache)
 - Good if more writes to the location follow
 - No-write-allocate (writes straight to memory, does not load into cache)
- **Typical**
 - Write-through + No-write-allocate
 - Write-back + Write-allocate

Intel Core i7 Cache Hierarchy

Processor package



L1 i-cache and d-cache:

32 KB, 8-way,

Access: 4 cycles

L2 unified cache:

256 KB, 8-way,

Access: 10 cycles

L3 unified cache:

8 MB, 16-way,

Access: 40-75 cycles

Block size: 64 bytes
for all caches.

Cache Performance Metrics

■ Miss Rate

- Fraction of memory references not found in cache (misses / accesses)
= $1 - \text{hit rate}$
- Typical numbers (in percentages):
 - 3-10% for L1
 - can be quite small (e.g., < 1%) for L2, depending on size, etc.

■ Hit Time

- Time to deliver a line in the cache to the processor
 - includes time to determine whether the line is in the cache
- Typical numbers:
 - 4 clock cycle for L1
 - 10 clock cycles for L2

■ Miss Penalty

- Additional time required because of a miss
 - typically 50-200 cycles for main memory (Trend: increasing!)

Cache Performance Metrics

- **Average memory access time (AMAT)**
 - Each request may hit or miss.
 - What is the average time over many requests?
- **AMAT = H + MR * AMP**
 - H = hit latency
 - MR = miss rate
 - AMP = average miss penalty

Cache Performance Metrics

- $\text{AMAT} = H + MR * \text{AMP}$
 - H = hit latency
 - MR = miss rate
 - AMP = average miss penalty

- Suppose we're given:

- Hit latency of L1 = 1 cycle
- Hit rate of L1 = 30%
- Latency of main memory = 200 cycles

- What is AMAT?

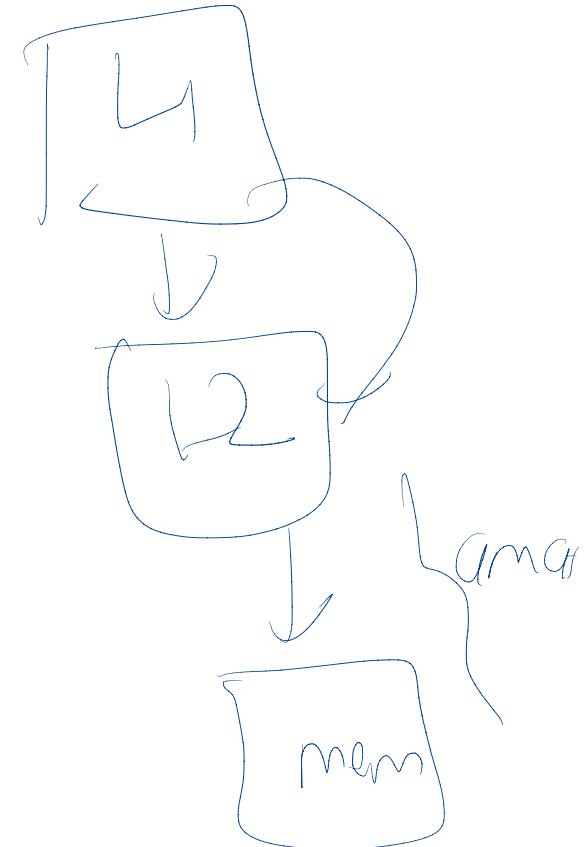
$$1 + .7 * 200 = 141$$

1 cycle to see if we have a hit or not

Misses are 70%

Cache Performance Metrics

- **AMAT = H + MR * AMP**
 - H = hit latency
 - MR = miss rate
 - AMP = average miss penalty
- **Suppose we're given:**
 - Hit latency of L1 = 1 cycle
 - Hit rate of L1 = 30%
 - Latency of main memory = 200 cycles
- **What is AMAT?**
 - $1 + 0.70 * 200 = 141$ cycles



Cache Performance Metrics

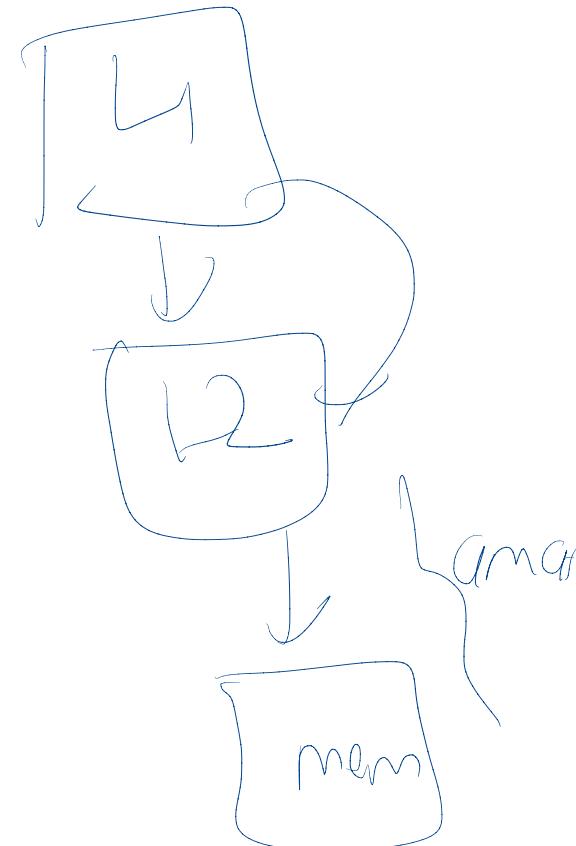
- What about multi-level caches?
 - $AMAT(L1) = H(L1) + MR(L1) * AMAT(L2)$
 - $AMAT(L2) = H(L2) + MR(L2) * AMP$

- Suppose we're given:

- L1: 1 cycle, 30% hit rate $1 + .7(7)$
- L2: 10 cycles, 55% hit rate $10 + .45 \times 200$
- RAM: 200 cycles

- What is AMAT?

$$5 + .3(10 + .2 \times 100)$$



Cache Performance Metrics

- What about multi-level caches?
 - $\text{AMAT(L1)} = H(L1) + MR(L1) * \text{AMAT}(L2)$
 - $\text{AMAT}(L2) = H(L2) + MR(L2) * \text{AMP}$
- Suppose we're given:
 - L1: 1 cycle, 30% hit rate
 - L2: 10 cycles, 55% hit rate
 - RAM: 200 cycles
- What is AMAT?
 - $1 + 0.70 * \text{AMAT}(L2)$

Cache Performance Metrics

■ What about multi-level caches?

- $AMAT(L1) = H(L1) + MR(L1) * AMAT(L2)$
- $AMAT(L2) = H(L2) + MR(L2) * AMP$

■ Suppose we're given:

- L1: 1 cycle, 30% hit rate
- L2: 10 cycles, 55% hit rate
- RAM: 200 cycles

$$S + 0.3(10 + 0.2 \cancel{+ 200})$$

■ What is AMAT?

- $1 + 0.70 * (10 + 0.45 * 200) = 71 \text{ cycles}$

Let's think about those numbers

- Huge difference between a hit and a miss
 - Could be 100x, if just L1 and main memory
- Would you believe 99% hits is twice as good as 97%?
 - Consider:
cache hit time of 1 cycle
miss penalty of 100 cycles
 - Average access time:
97% hits: 1 cycle + 0.03 * 100 cycles = **4 cycles**
99% hits: 1 cycle + 0.01 * 100 cycles = **2 cycles**
factor of L
- This is why “miss rate” is used instead of “hit rate”

Writing Cache Friendly Code

- **Make the common case go fast**
 - Focus on the inner loops of the core functions
- **Minimize the misses in the inner loops**
 - Repeated references to variables are good (**temporal locality**)
 - Stride-1 reference patterns are good (**spatial locality**)

Key idea: Our qualitative notion of locality is quantified through our understanding of cache memories

Cache Performance

■ What is the hit rate of this code?

- 64 B block
- int is 4 B

```
int a[m][n];
for (int i = 0; i < m; i++)
    for (int j = 0; j < n; j++)
        a[i][j] += 1;
```

Cache Performance

■ What is the hit rate of this code?

- 64 B block
- Int is 4 B

```
int a[m][n];
for (int i = 0; i < m; i++)
    for (int j = 0; j < n; j++)
        a[i][j] += 1;
```

- For each cache line (64 B), it holds 16 ints
- First will miss, next 15 will hit
- Then another miss, then another 15 hits, ...
- Hit rate = 15/16

Matrix Multiplication Example

■ Description:

- Multiply $N \times N$ matrices
- Matrix elements are doubles (8 bytes)
- $O(N^3)$ total operations
- N reads per source element
- N values summed per destination
 - but may be able to hold in register

```
/* ijk */  
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        sum = 0.0; ←  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum;  
    }  
}
```

*Variable **sum**
held in register*

matmult/mm.c

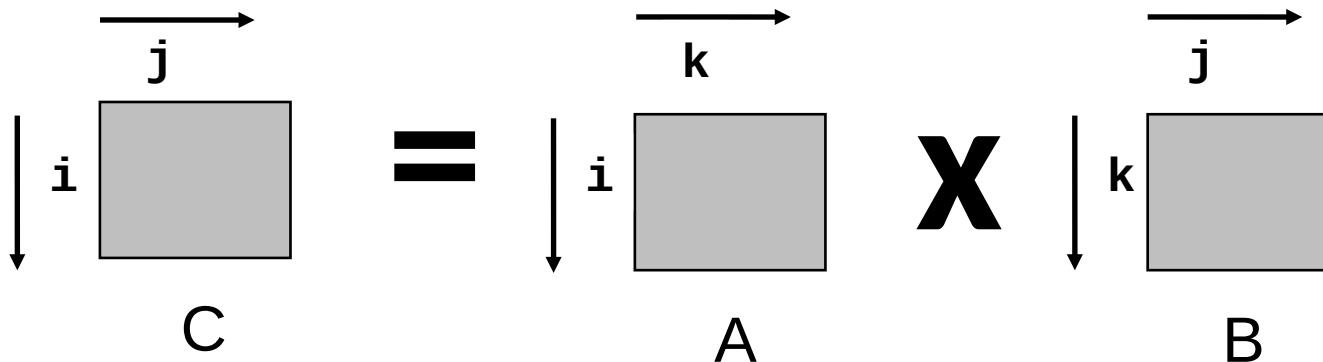
Miss Rate Analysis for Matrix Multiply

■ Assume:

- Block size = $32B$ (big enough for four doubles)
- Matrix dimension (N) is very large
 - Approximate $1/N$ as 0.0
- Cache is not even big enough to hold multiple rows

■ Analysis Method:

- Look at access pattern of inner loop



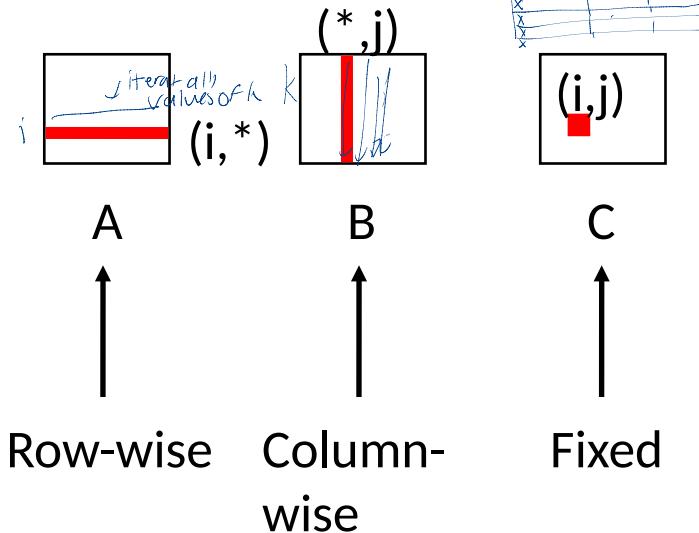
Layout of C Arrays in Memory (review)

- C arrays allocated in row-major order
 - each row in contiguous memory locations
- Stepping through columns in one row:
 - ```
for (i = 0; i < N; i++)
 sum += a[0][i];
```
  - accesses successive elements
  - if block size (B) > sizeof( $a_{ij}$ ) bytes, exploit spatial locality
    - miss rate =  $\text{sizeof}(a_{ij}) / B$
- Stepping through rows in one column:
  - ```
for (i = 0; i < n; i++)
    sum += a[i][0];
```
 - accesses distant elements
 - no spatial locality!
 - miss rate = 1 (i.e. 100%)

Matrix Multiplication (ijk)

```
/* ijk */  
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum;  
    }  
}  
  
matmult/mm.c
```

Inner loop:



Misses per inner loop iteration:

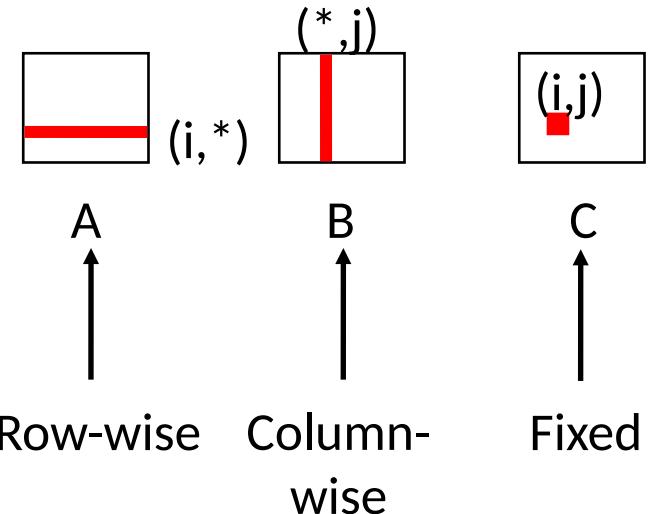
A	B	C
0.25	1.0	0.0

Matrix Multiplication (jik)

```
/* jik */  
for (j=0; j<n; j++) {  
    for (i=0; i<n; i++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum  
    }  
}
```

matmult/mm.c

Inner loop:



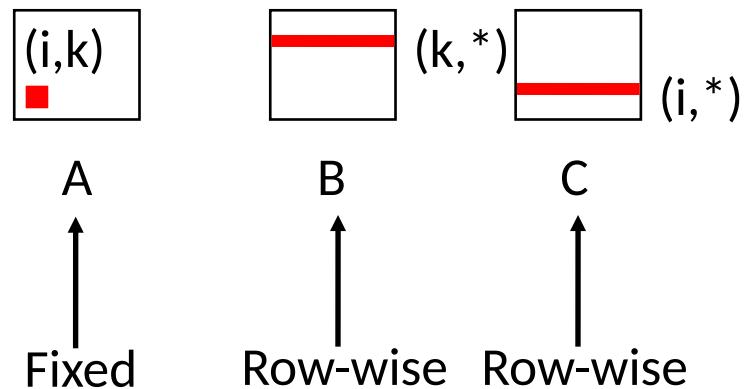
Misses per inner loop iteration:

A	B	C
0.25	1.0	0.0

Matrix Multiplication (kij)

```
/* kij */  
for (k=0; k<n; k++) {  
    for (i=0; i<n; i++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}  
matmult/mm.c
```

Inner loop:



Misses per inner loop iteration:

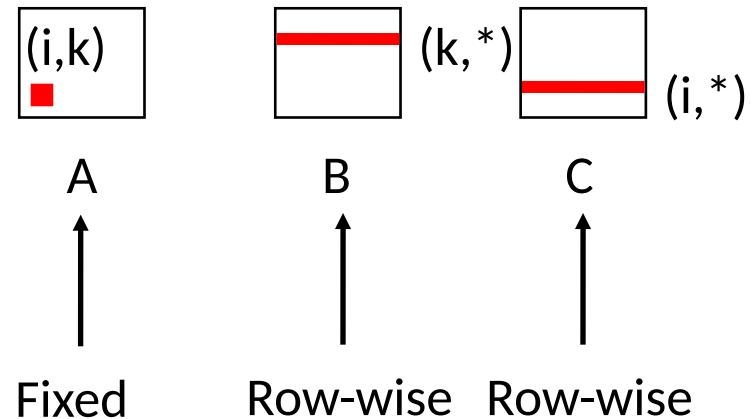
A	B	C
0.0	0.25	0.25

Matrix Multiplication (ikj)

```
/* ikj */  
for (i=0; i<n; i++) {  
    for (k=0; k<n; k++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```

matmult/mm.c

Inner loop:



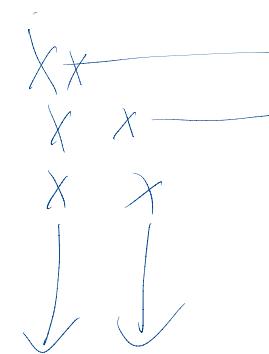
Misses per inner loop iteration:

A	B	C
0.0	0.25	0.25

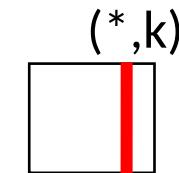
Matrix Multiplication (jki)

```
/* jki */  
for (j=0; j<n; j++) {  
    for (k=0; k<n; k++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```

matmult/mm.c



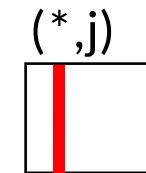
Inner loop:



A



B



C

Column-wise
↑
Fixed

Column-wise
↑
Column-wise

Misses per inner loop iteration:

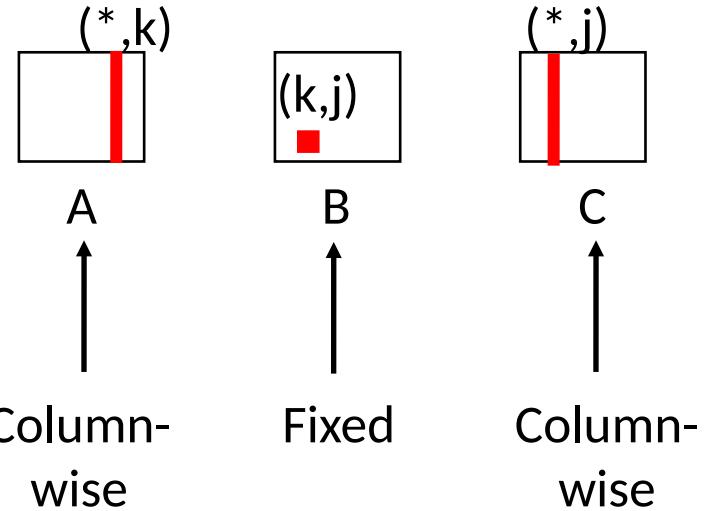
A	B	C
1.0	0.0	1.0

Matrix Multiplication (kji)

```
/* kji */
for (k=0; k<n; k++) {
    for (j=0; j<n; j++) {
        r = b[k][j];
        for (i=0; i<n; i++)
            c[i][j] += a[i][k] * r;
    }
}
```

matmult/mm.c

Inner loop:



Misses per inner loop iteration:

A	B	C
1.0	0.0	1.0

Summary of Matrix Multiplication

```
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum;  
    }  
}
```

```
for (k=0; k<n; k++) {  
    for (i=0; i<n; i++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```

```
for (j=0; j<n; j++) {  
    for (k=0; k<n; k++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```

ijk (& jik):

- 2 loads, 0 stores
- misses/iter = **1.25**

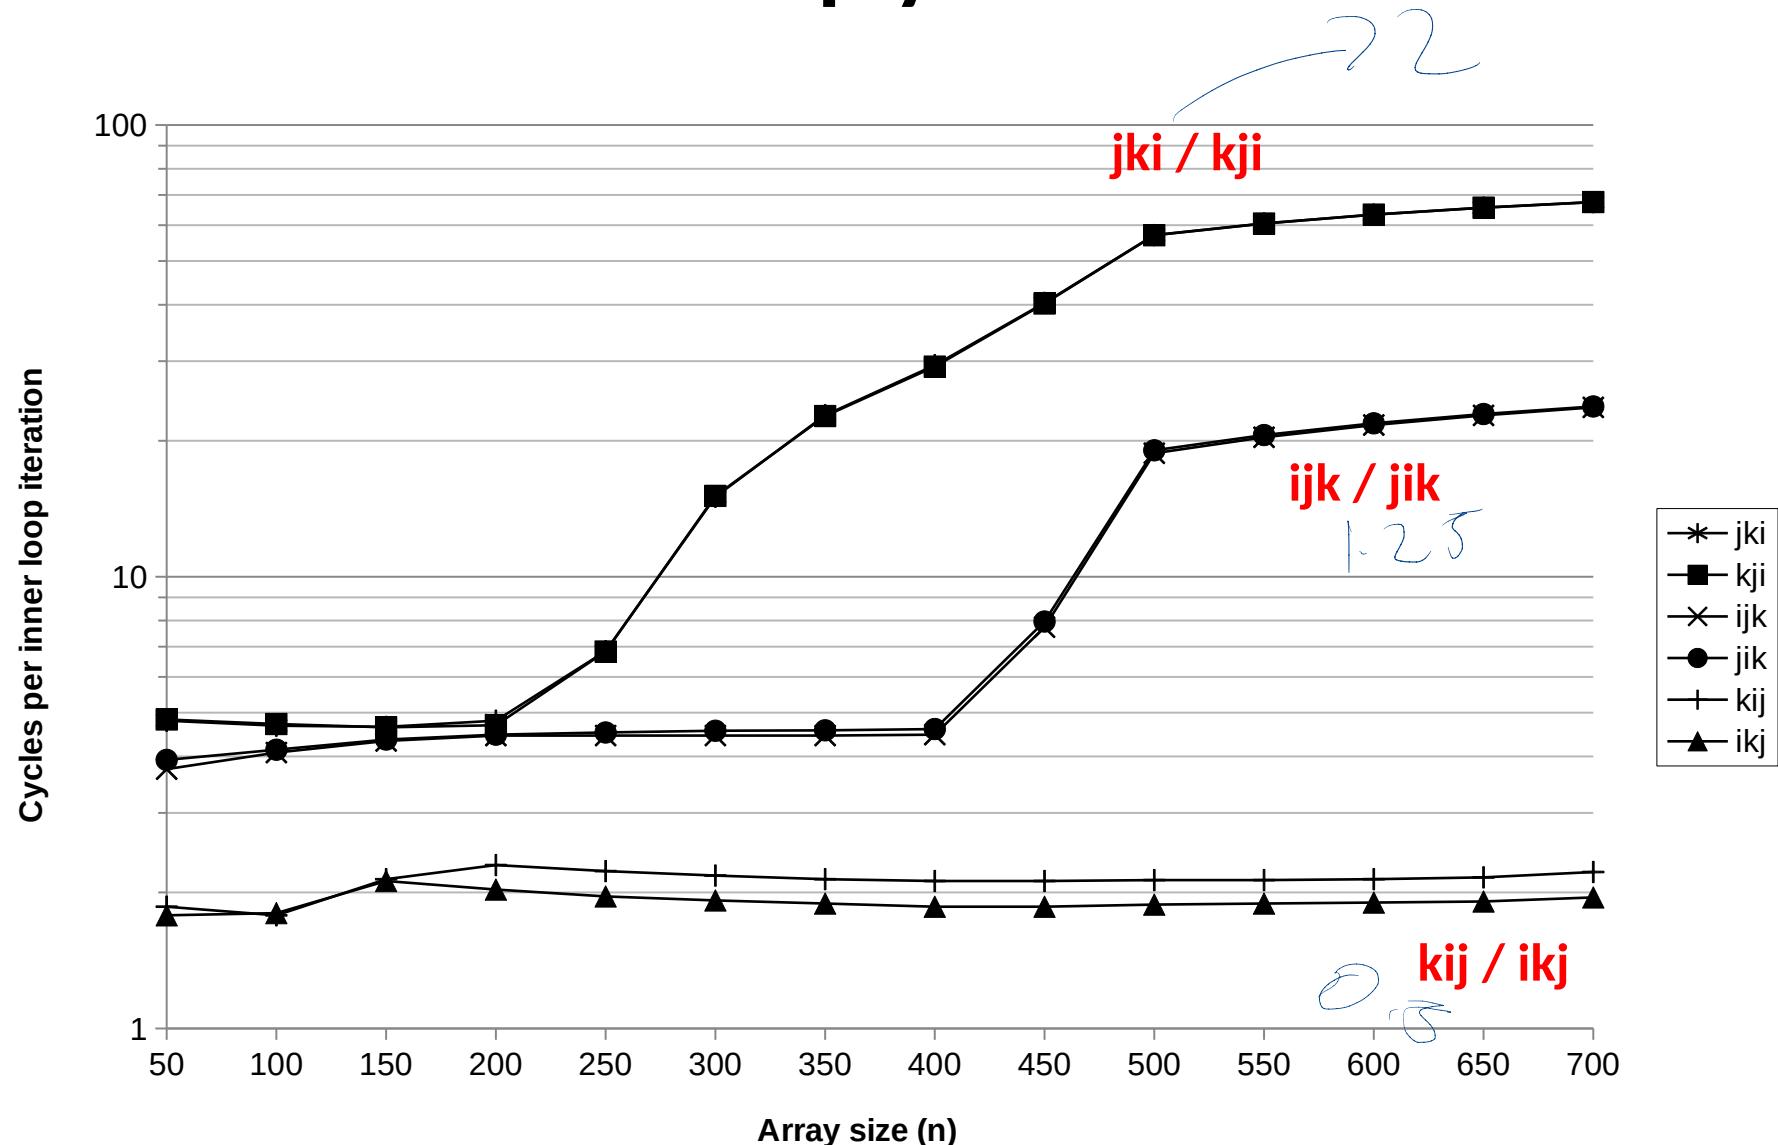
kij (& ikj):

- 2 loads, 1 store
- misses/iter = **0.5**

jki (& kji):

- 2 loads, 1 store
- misses/iter = **2.0**

Core i7 Matrix Multiply Performance



Cache Summary

- Cache memories can have significant performance impact
- You can write your programs to exploit this!
 - Focus on the inner loops, where bulk of computations and memory accesses occur.
 - Try to maximize spatial locality by reading data objects with sequentially with stride 1.
 - Try to maximize temporal locality by using a data object as often as possible once it's read from memory.